Logic and Mechanized Reasoning

Release 0.1

Jeremy Avigad
Seul Baek
Marijn J. H. Heule
Wojciech Nawrocki
Emre Yolcu

Oct 11, 2021

CONTENTS:

INTRODUCTION

1.1 Historical background

In the thirteenth century, Ramon Lull, an eccentric Franciscan monk on the Island of Majorca, wrote a work called the *Ars Magna*, which contains mechanical devices designed to aid reasoning. For example, Lull claimed that the reason that infidels did not accept the Christian god was that they failed to appreciate the multiplicity of God's attributes, so he designed nested paper circles with letters and words signifying these attributes. By rotating the circles, one obtained various combinations of the words and symbols, signifying compound attributes. For example, in one configuration, one might read off the result that God's greatness is true and good, and that his power is wise and great. Other devices and diagrams would assist in reasoning about virtues and vices, and so on.

Today, this sounds silly, but the work was based on three fundamental ideas that are still of central importance.

- First, we can use symbols, or tokens, to stand for ideas or concepts.

- Second, compound ideas and concepts are formed by putting together simpler ones.

- And, third, mechanical devices—even as simple a concentric rotating wheels—can serve as aids to reasoning.

The first two ideas go back to ancient times and can be found in the work of Aristotle. The third, however, is usually attributed to Lull, marking his work as the origin of mechanized reasoning.

Four centuries later, Lull's work resonated with Gottfied Leibniz, who invented calculus around the same time that Isaac Newton did so independently. Leibniz was also impressed by the possibility of symbolic representations of concepts and rules for reasoning. He spoke of developing a *characteristica universalis*—a universal language of thought—and a *calculus ratiocinator*—a calculus for reasoning. In a famous passage, he wrote:

> If controversies were to arise, there would be no more need of disputation between two philosophers than between two calculators. For it would suffice for them to take their pencils in their hands and to sit down at the abacus, and say to each other (and if they so wish also to a friend called to help): Let us calculate.

The last phrase—*calculemus!* in the original Latin—has become a motto of computer scientists and computationally-minded mathematicians today.

The development of modern logic in the late eighteenth and early nineteen centuries began to bring Leibniz' vision to fruition. In 1931, Kurt Gödel wrote:

> The development of mathematics towards greater precision has led, as is well known, to the formalization of large tracts of it, so that one can prove any theorem using nothing but a few mechanical rules.

It is notable that the use of the word "mechanical" here—*mechanischen* in the original German—predates the modern computer by a decade or so.

What logicians from the time of Aristotle to the present day have in common is that they are all at least slightly crazy. They are driven by the view that knowledge is rooted in language and that the key to knowledge lies in having just the right symbolic representations of language and rules of use. But often it's the crazy people that change the world. The logical view of language and knowledge lies at the heart of computer science and provides the foundation for some of our most valuable technologies today, including programming languages, automated reasoning and AI, and databases.

That's what this course is about: the logician's view of the world, the power of symbolic representations of language, and the way those representations facilitate the mechanization of reasoning and the acquisition of knowledge.

The logicians' view complements the view from statistics and machine learning, where representations of knowledge tend to be very large, approximate, and hard to represent is succinct symbolic terms. Such methods have had stunning successes in recent years, but there are still branches of computer science and AI where symbolic methods are paramount. It is an important open question as to the best way to combine logical, statistical, and machine learning methods in the years to come.

1.2 An overview of this course

This course is designed to teach you the mathematical theory behind symbolic logic, with an eye towards putting it to good use. An interesting aspect of the course is that it develops three interacting strands in parallel:

- *Theory.* We will teach you the syntax and semantics of propositional and first-order logic. If time allows, we will give you a brief overview of related topics, like simple type theory and higher-order logic.

- *Implementation.* We will teach you how to implement logical syntax—terms and formulas—in a functional programming language called *Lean*. We will also teach you how to carry out fundamental operations and transformations on these objects.

- *Application.* We will show you how to use logic-based automated reasoning tools to solve interesting and difficult problems. In particular, we will use a SAT solver called CaDiCaL, an SMT solver called Z3, and a first-order theorem prover called Vampire (and by then you will understand what all these terms mean).

The first strand will be an instance of pure mathematics. We will build on the skills you have learned in Mathematical Foundations of Computer Science (15-151). The goal is to teach you to think about and talk about logic in a mathematically rigorous way.

The second strand will give you an opportunity to code up some of what you have learned and put it to good use. Our goal is to provide a foundation for you to use logic-based computational methods in the future, whether you choose to make use of them in small or large ways. In the third strand, for illustrative purposes, we will focus mainly on solving puzzles and combinatorial problems. This will give you a sense of how the tools can also be used on proof and constraint satisfaction problems that come up in fields like program verification, discrete optimization, and AI.

MATHEMATICAL BACKGROUND

2.1 Induction and recursion on the natural numbers

In its most basic form, the principle of induction on the natural numbers says that if you want to prove that every natural number has some property, it suffices to show that zero has the property, and that whenever some number n has the property, so does $n + 1$. Here is an example.

Theorem

For every natural number n, $\sum_{i \leq n} i = n(n+1)/2$.

Proof

Use induction on n. In the base case, we have $\sum_{i \leq 0} i = 0 = 0(0+1)/2$. For the induction step, assuming $\sum_{i \leq n} i = n(n+1)/2$, we have

$$\sum_{i \leq n+1} i = \sum_{i \leq n} i + (n+1)$$
$$= n(n+1)/2 + 2(n+1)/2$$
$$= (n+1)(n+2)/2$$

The story is often told that Gauss, as a schoolchild, discovered this formula by writing

$$S = 1 + \ldots + n$$
$$S = n + \ldots + 1$$

and then adding the two rows and dividing by two. The proof by induction doesn't provide insight as to how one might *discover* the theorem, but once you have guessed it, it provides a short and effective means for establishing that it is true.

In a similar vein, you might notice that an initial segment of the odd numbers yields a perfect square. For example, we have $1 + 3 + 5 + 7 + 9 = 25$. Here is a proof of the general fact:

Theorem

For every natural number n, $\sum_{i \leq n}(2i + 1) = (n + 1)^2$.

Proof

The base case is easy, and assuming the inductive hypothesis, we have

$$\sum_{i \leq n+1} (2i+1) = \sum_{i \leq n} (2i+1) + 2(n+1) + 1$$
$$= (n+1)^2 + 2n + 3$$
$$= n^2 + 4n + 4$$
$$= (n+2)^2.$$

A close companion to induction is the principle of *recursion*. Recursion enables us to define functions on the natural numbers, and induction allows us to prove things about them. For example, let $g : \mathbb{N} \to \mathbb{N}$ be the function defined by

$$g(0) = 1$$
$$g(n+1) = (n+1) \cdot g(n)$$

Then g is what is known as the *factorial* function, whereby $g(n)$ is conventionally written $n!$. The point is that if you don't know what the factorial function is, the two equations above provide a complete specification. There is exactly one function, defined on the natural numbers, that meets that description.

Here is an identity involving the factorial function:

Theorem

$\sum_{i \leq n} i \cdot i! = (n+1)! - 1$.

Proof

The base case is easy. Assuming the claim holds for n, we have

$$\sum_{i \leq n+1} i \cdot i! = \sum_{i \leq n} i \cdot i! + (n+1) \cdot (n+1)!$$
$$= (n+1)! + (n+1) \cdot (n+1)! - 1$$
$$= (n+1)!(1 + (n+1)) - 1$$
$$= (n+2)! - 1$$

This is a pattern found throughout mathematics and computer science: define functions and operations using recursion, and then use induction to prove things about them.

The *Towers of Hanoi* puzzle provides a textbook example of a problem that can be solved recursively. The puzzle consists of three pegs and disks of different diameters that slide onto the pegs. The initial configuration has n disks stacked on one of the pegs in decreasing order, with the largest one at the bottom and the smallest one at the top. Suppose the pegs are numbered 1, 2, and 3, with the disks starting on peg 1. The required task is to move all the disks from peg 1 to peg 2, one at a time, with the constraint that a larger disk is never placed on top of a smaller one.

```
To move n disks from peg A to peg B with auxiliary peg C:
    if n = 0
        return
    else
        move n − 1 disks from peg A to peg C using auxiliary peg B
        move 1 disk from peg A to peg B
        move n − 1 disks from peg C to peg B using auxiliary peg A
```

We will show in class that this requires $2^n - 1$ moves. The exercises below ask you to show that *any* solution requires at least this many moves.

2.2 Complete induction

As we have described it, the principle of induction is pretty rigid: in the inductive step, to show that $n + 1$ has some property, we can only use the corresponding property of n. The principle of *complete* induction is much more flexible.

Principle of complete induction

To show that every natural number n has some property, show that n has that property whenever all smaller numbers do.

As an exercise, we ask you to prove the principle of complete induction using the ordinary principle of induction. Remember that a natural number greater than or equal to 2 is *composite* if it can be written as a product of two smaller numbers, and *prime* otherwise.

Theorem

Every number greater than two can be factored into primes.

Proof

Let n be any natural number greater than or equal to 2. If n is prime, we are done. Otherwise, write $n = m \cdot k$, where m and k are smaller than n (and hence greater than 1). By the inductive hypothesis, m and k can each be factored into prime numbers, and combining these yields a factorization of n.

Here is another example we will discuss in class:

Theorem

For any $n \geq 3$, the sum of the angles in any n-gon is $180(n - 2)$.

The companion to complete induction on the natural numbers is a form of recursion known as course-of-values recursion, which allows you to define a function f by giving the value of $f(n)$ in terms of the value of f at arbitrary smaller values of n. For example, we can define the sequence of Fibonacci numbers as follows:

$$F_0 = 0$$
$$F_1 = 1$$
$$F_{n+2} = F_{n+1} + F_n$$

The fibonacci numbers satisfy lots of interesting identities, some of which are given in the exercises.

In fact, you can define a function by recursion as long as *some* associated measure decreases with each recursive call. Define a function $f(n, k)$ for $k \leq n$ by

$$f(n, k) = \begin{cases} 1 & \text{if } k = 0 \text{ or } k = n \\ f(n-1, k) + f(n-1, k-1) & \text{otherwise} \end{cases}$$

Here it is the first argument that decreases. In class, we'll discuss a proof that this defines the function

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

which is simultaneously equal to number of ways of choosing k objects out of n without repetition.

Finally, here is a recursive description of the greatest common divisor of two nonnegative integers:

$$gcd(x, y) = \begin{cases} x & \text{if } y = 0 \\ gcd(y, mod(x, y)) & \text{otherwise} \end{cases}$$

where $mod(x, y)$ is the remainder when dividing x by y.

2.3 Generalized induction and recursion

The natural numbers are characterized inductively by the following clauses:

- 0 is a natural number.

- If x is a natural number, so is $succ(x)$.

Here the function $succ(x)$ is known as the *successor* function, namely, the function that, given any number, returns the next one in the sequence. The natural numbers structure is also sometimes said to be *freely generated* by this data. The fact that it is *generated* by 0 and $succ(x)$ means that it is the *smallest* set that contains 0 and is closed under $succ(x)$; in other words, any set of natural numbers that contains 0 and is closed under $succ(x)$ contains all of them. This is just the principle of induction in disguise. The fact that it is generated *freely* by these elements means that there is no confusion between them: 0 is not a successor, and if $succ(x) = succ(y)$, then $x = y$. Intuitively, being generated by 0 and $succ(x)$ means that any number can be represented by an expression built up from these, and being generated freely means that the representation is unique.

The natural numbers are an example of an *inductively defined structure*. These come up often in logic and computer science. It is often useful to define functions by recursion on such structures, and to use induction to prove things about them. We will describe the general schema here with some examples that often come up in computer science.

Let α be any data type. The set of all *lists* of elements of α, which we will write as $List(\alpha)$, is defined inductively as follows:

- The element *nil* is an element of $List(\alpha)$.

- If a is an element of α and ℓ is an element of $List(\alpha)$, then the element $cons(a, \ell)$ is an element of $List(\alpha)$.

Here *nil* is intended to describe the empty list, $[]$, and $cons(a, \ell)$ is intended to describe the result of adding a to the beginning of ℓ. So, for example, the list of natural numbers $[1, 2, 3]$ would be written $cons(1, cons(2, cons(3, nil)))$. Think of $List(\alpha)$ as having a constructor $cons(a, \cdot)$ for each a. Then, in the terminology above, $List(\alpha)$ is generated inductively by *nil* and those constructors.

Henceforth, for clarity, we'll use the notation $[]$ for *nil* and $a :: \ell$ for $cons(a, \ell)$. More generally, we can take $[a, b, c, \ldots]$ to be an abbreviation for $a :: (b :: (c :: \ldots []))$.

Saying that $List(\alpha)$ is inductively defined means that we principles of recursion and induction on it. For example, the following concatenates two lists:

$$append([], m) = m$$
$$append((a :: \ell), m) = a :: (append(\ell, m))$$

Here the recursion is on the first argument. As with the natural numbers, the recursive definition specifies what to do for each of the constructors. We'll use the notation $\ell \mathbin{+\!\!+} m$ for $append(\ell, m)$, and with this notation, the two defining clauses read as follows:

$$[] \mathbin{+\!\!+} m = m$$
$$(a :: \ell) \mathbin{+\!\!+} m = a :: (\ell \mathbin{+\!\!+} m)$$

From the definition, we have $[] \mathbin{+\!\!+} \ell = \ell$ for every ℓ, but $m \mathbin{+\!\!+} [] = m$ is something we have to prove.

Proposition

For every m, we have $m + [] = m$.

Proof

We use induction on m. In the base case, we have $[] + [] = []$ from the definition of *append*. For the induction step, suppose we have $m + [] = m$. Then we also have

$$(a :: m) + [] = a :: (m + [])$$
$$= a :: m.$$

The definition of the append function is an example of *structural recursion*, called that because the definition proceeds by recursion on the structure of the inductively defined type. In particular, there is a clause of the definition corresponding to each constructor. The proof we have just seen is an instance of *structural induction*, called that because, once again, there is part of the proof for each constructor. The base case, for *nil*, is straightforward, because that constructor has no arguments. The inductive step, for *cons*, comes with an inductive hypothesis because the *cons* constructor has a recursive argument. In class, we'll do a similar proof that the *append* operation is associative.

The following function (sometimes called *snoc*) appends a single element at the end:

$$append1([], a) = cons(a, nil)$$
$$append1(cons(b, \ell), a) = cons(b, append1(\ell, a))$$

An easy induction on ℓ shows that, as you would expect, $append1(\ell, a)$ is equal to $\ell + [a]$.

The following function reverses a list:

$$reverse([]) = []$$
$$reverse(cons(a, \ell)) = append1(reverse(\ell), a)$$

In class, or for homework, we'll work through proofs that that the following holds for every pair of lists ℓ and m:

$$reverse(\ell + m) = reverse(m) + reverse(\ell)$$

Here is another example of a property that can be proved by induction:

$$reverse(reverse(\ell)) = \ell$$

From a mathematical point of view, this definition of the reverse function above is as good as any other, since it specifies the function we want unambiguously. But in Chapter 3 we will see that such a definition can also be interpreted as executable code in a functional programming language such as Lean. In this case, the execution is quadratic in the length of the list (think about why). The following definition is more efficient in that sense:

$$reverseAux([], m) = m$$
$$reverseAux((a :: \ell), m) = reverseAux(\ell, (a :: m))$$

$$reverse'(\ell) = reverseAux(\ell, [])$$

The idea is that *reverseAux* adds all the elements of the first argument to the second one in reverse order. So the second arguments acts as an *accumulator*. In fact, because it is a tail recursive description, the code generated by Lean is quite efficient. In class, we'll discuss an inductive proof that $reverse(\ell) = reverse'(\ell)$ for every ℓ.

It is worth mentioning that structural induction is not the only way to prove things about lists, and structural recursion is not the only way to define functions by recursion. Generally speaking, we can assign any complexity measure to a data type, and do induction on complexity, as long as the measure is well founded. (This will be the case, for example, for measures that take values in the natural numbers, with the usual ordering on size.) For example, we can define the length of a list as follows:

$$length([]) = 0$$
$$length(a :: \ell) = length(\ell) + 1$$

Then we can define a function f on lists by giving the value of $f(\ell)$ in terms of the value of f on smaller lists, and we can prove a property of lists using the fact that the property holds of all smaller lists as an inductive hypothesis. These are ordinary instances of recursion and induction on the natural numbers.

As another example, we consider the type of finite binary trees, defined inductively as follows:

- The element *empty* is a binary tree.

- If s and t are finite binary trees, so is the $node(s, t)$.

In this definition, *empty* is intended to denote the empty tree, and $node(s, t)$ is intended to denote the binary tree that consists of a node at the top and has s and t as the left and right subtrees, respectively.

Be careful: it is more common to take the set of binary trees to consist of only the *nonempty* trees, in which case, what we have defined here are called the *extended* binary trees. Adding the empty tree results in a nice inductive characterization. If we started with a one-node tree as the base case, we would have to allow for three types of compound tree: one type with a node and a subtree to the left, one with a node and a subtree to the right, and one with a node with both left and right subtrees.

We can count the number of nodes in an extended binary tree with the following recursive definition:

$$size(empty) = 0$$
$$size(node(s, t)) = 1 + size(s) + size(t)$$

We can compute the depth of an extended binary tree as follows:

$$depth(empty) = 0$$
$$depth(node(s, t)) = 1 + \max(depth(s), depth(t))$$

Again, be careful: many authors take the depth of a tree to be the length of the longest path from the root to a leaf, in which case, what we have defined here computes the depth *plus one* for nonempty trees.

2.4 Invariants

The *mutilated chessboard* problems involves an 8×8 chessboard with the top right and bottom left corners removed. Imagine you are given a set of dominoes, each of which can cover exactly two squares. It is possible to cover all the squares of the mutilated chessboard using dominoes, so that each square is covered by exactly one domino?

A moment's reflection shows that the answer is no. If you imagine the chessboard squares colored white and black in the usual way, you'll notice that the two squares we removed have the same color, say, black. That means that there are more white squares than black squares. On the other hand, every domino covers exactly one square of each color. So no matter how many dominoes we put down, we'll never have them color more white squares than black squares.

The fact that any way of putting down dominoes covers the same number of white and black squares is an instance of an *invariant*, which is a powerful idea in both mathematics and computer science. An invariant is something—a quantity, or a property—that doesn't change as something else does (in this case, the number of dominoes).

Often the natural way to establish an invariant uses induction. In this case, it is obvious that putting down one domino doesn't change the difference between the number of white and black squares covered, since each domino covers one of each. By induction on n, putting down n dominoes doesn't change the difference either.

The following puzzle, called the *MU puzzle*, comes from the book *Gödel, Escher, Bach* by Douglas Hofstadter. It concerns strings consisting of the letters *M*, *I*, and *U*. Starting with the string *MI*, we are allowed to apply any of the following rules:

1. Replace *sI* by *sIU*, that is, add a *U* to the end of any string that ends with *I*.

2. Replace *Ms* by *Mss*, that is, double the string after the initial *M*.

3. Replace *sIIIt* by *sUt*, that is, replace any three consecutive *Is* with a *U*.

4. Replace *sUUt* by *st*, that is, delete any consecutive pair of *Us*.

The puzzle asks whether it is possible to derive the string *MU*. The answer is no: it turns out that a string is derivable if and only if it consists of an *M* followed by any number of *Is* and *Us*, as long as the number of *Is* is not divisible by 3. In class, we'll prove the "only if" part of this equivalence. Try the "if" part if you like a challenge.

As a final example, in class we'll discuss the Golomb *tromino theorem*. A *tromino* is an L-shaped configuration of three squares. Golomb's theorem says that any $2^n \times 2^n$ chessboard with one square removed can be tiled with trominoes. We'll prove this together in class.

2.5 Exercises

1. Prove the formula for the sum of a geometric series:

$$\sum_{i<n} ar^i = \frac{a(r^n - 1)}{r - 1}$$

2. Prove that for every $n > 4$, $n! > 2^n$.

3. Show that the solution to the towers of Hanoi given in Section 2.1 is optimal: for every n, it takes at least *2^n - 1* moves to move all the disks from one peg to another.

4. Consider the variation on the towers of Hanoi problem in which you can only move a disk to an *adjacent* peg. In other words, you can move a disk from peg 1 to peg 2, from peg 2 to peg 1, from peg 2 to peg 3, or from peg 3 to peg 2, but not from peg 1 to peg 3 or from peg 3 to peg 1.

Describe a recursive procedure for solving this problem, and show that is requires $3^n - 1$ moves. If you are ambitious, show that this is optimal, and that it goes through all the 3^n valid positions.

5. Consider the variation on the towers of Hanoi in which pegs can be moved cyclicly: from peg 1 to peg 2, from peg 2 to peg 3, or from peg 3 to peg 1. Describe a recursive procedure to solve this problem.

6. Use the ordinary principle of induction to prove the principle of complete induction.

7. Let $F_0, F_1, F_2, \ldots$ be the sequence of Fibonacci numbers.

   1. Let α and β be the two roots of the equation $x^2 = x + 1$. Show that for every n, $F_n = (\alpha^n - \beta^n)/\sqrt{5}$.

   2. Show $\sum_{i<n} F_i = F_{n+1} - 1$.

   3. Show $\sum_{i\leq n} F_i^2 = F_n F_{n+1}$.

8. Show that with n straight lines we can divide the plane into at most $n^2 + n + 2$ regions, and that this is sharp.

9. Show that the recursive description of $gcd(x, y)$ presented in Section 2.3 correctly computes the greatest common divisor of x and y, where we define $gcd(0, 0)$ equal to 0. You can restrict attention to nonnegative values of x and y. (Hint: you can use the fact that for for every y not equal to 0, we can write $x = div(x, y) \cdot y + mod(x, y)$, where $div(x, y)$ is the integer part of x divided by y. First show that for every k, $gcd(x, y) = gcd(x + ky, y)$, and use that fact.

10. Use structural induction to prove

$$reverse(\ell + m) = reverse(m) + reverse(\ell).$$

11. Use structural induction to prove

$$reverse(reverse(\ell)) = \ell.$$

12. Prove that for every ℓ we have

$$reverse'(\ell) = reverse(\ell).$$

13. Prove that for every ℓ and m we have

$$length(\ell + m) = length(\ell) + length(m).$$

14. How many binary trees of depth n are there? Prove your answer is correct.

15. Show that a string is derivable in the MU puzzle if and only if it consists of an M followed by any number of Is and Us, as long as the number of Is is not divisible by 3.

USING LEAN AS A PROGRAMMING LANGUAGE

3.1 About Lean

Lean is a new programming language and interactive proof assistant being developed at Microsoft Research. It is currently in an experimental, development stage, which makes it a risky choice for this course. But in many ways it is an ideal system for working with logical syntax and putting logic to use. Lean is an exciting project, and the system fun to use. So please bear with us. Using Lean puts us out on the frontier, but if you adopt a pioneering attitude, you will be in a good position to enjoy all the cool things that Lean has to offer.

You can learn more about Lean on the Lean home page, on the Lean community home page, and by asking questions on the Lean Zulip chat, which you are heartily encouraged to join. To be more precise, there are currently two versions of Lean:

- Lean 3 is reasonably stable, and primarily an interactive proof assistant. It has a very large mathematical library, known as mathlib.

- Lean 4 is being designed as a performant programming language, and it is still under development. It can also be used as a proof assistant, though it does not yet have a substantial library. Its language and syntax are similar to that of Lean 3, but it is not backward compatible.

In this course, we will use Lean 4, even though it is still under development. It has the rough beginnings of a user manual and there is a tutorial on the underlying foundation. As we will see, Lean has a lot of features that make that worthwhile. In particular, Lean 4 is designed to be an ideal language for implementing powerful logic-based systems, as evidenced by the fact that most of Lean 4 is implemented in Lean 4 itself.

The goal of this section is to give you a better sense of what Lean is, how it can possibly be a programming language and proof assistant at the same time, and why that makes sense. The rest of the introduction will give you a quick tour of some of its features, and we will learn more about them as the course progresses.

At the core, Lean is an implementation of a formal logical foundation known as *type theory*. More specifically, it is an implementation of *dependent type theory*, and even more specifically than that, it implements a version of the *Calculus of Inductive Constructions*. Saying that it implements a formal logic foundation means that there is a precise grammar for writing expressions, and precise rules for using them. In Lean, every well-formed expression has a type.

```
#check 2 + 2
#check -5
#check [1, 2, 3]
#check #[1, 2, 3]
#check (1, 2, 3)
#check "hello world"
#check true
#check fun x => x + 1
#check fun x => if x = 1 then "yes" else "no"
```

You can find this example in the file *using_lean_as_a_programming_language/examples1.lean* in the *LAMR/Examples* folder of the course repository. We recommend copying that entire folder in the *User* folder, so you can edit the files and try examples of your own. You can always find the original file in the folder *LAMR/Examples*, which you should not edit.

If you hover over the *#check* statements or move your cursor to one of these lines and check the information window, Lean reports the result of the command. It tells you that *2 + 2* has type *Nat*, *-5* has type *Int*, and so on. In fact, in the formal foundation, types are expressions as well. The types of all the expressions above are listed below:

```lean
#check Nat
#check Int
#check List Nat
#check Array Nat
#check Nat × Nat × Nat
#check String
#check Bool
#check Nat → Nat
#check Nat → String
```

Now Lean tells you each of these has type *Type*, indicating that they are all data types. If you know the type of an expression, you can ask Lean to confirm it:

```lean
#check (2 + 2 : Nat)
#check ([1, 2, 3] : List Nat)
```

Lean will report an error if it cannot construe the expression as having the indicated type.

In Lean, you can define new objects with the *def* command. The new definition becomes part of the *environment*: the defined expression is associated with the identifier that appears after the word *def*.

```lean
def four : Nat := 2 + 2

def isOne (x : Nat) : String := if x = 1 then "yes" else "no"

#check four
#print four

#check isOne
#print isOne
```

The type annotations indicate the intended types of the arguments and the result, but they can be omitted when Lean can infer them from the context:

```lean
def four' := 2 + 2

def isOne' x := if x = 1 then "yes" else "no"
```

So far, so good: in Lean, we can define expressions and check their types. What makes Lean into a programming language is that the logical foundation has a computational semantics, under which expressions can be *evaluated*.

```lean
#eval four
#eval isOne 3
#eval isOne 1
```

The *#eval* command evaluates the expression and then displays the return value. Evaluation can also have *side effects*, which are generally related to system IO. For example, displaying the string "Hello, world!" is a side effect of the following evaluation:

```
#eval IO.println "Hello, world!"
```

Theoretical computer scientists are used to thinking about programs as expressions and identifying the act of running the program with the act of evaluating the expression. In Lean, this view is made manifest, and the expressions are defined in a formal system with a precise specification.

But what makes Lean into a proof assistant? To start with, some expressions in the proof system express propositions:

```
#check 2 + 2 = 4
#check 2 + 2 < 5
#check isOne 3 = "no"
#check 2 + 2 < 5 ∧ isOne 3 = "no"
```

Lean confirms that each of these is a proposition by reporting that each of them has type *Prop*. Notice that they do not all express *true* propositions; theorem proving is about certifying the ones that are. But the language of Lean is flexible enough to express just about any meaningful mathematical statement at all. For example, here is the statement of Fermat's last theorem:

```
def Fermat_statement : Prop :=
∀ a b c n : Nat, a * b * c ≠ 0 ∧ n > 2 → a^n + b^n ≠ c^n
```

In Lean's formal system, data types are expressions of type *Type*, and if *T* is a type, an expression of type *T* denotes an object of that type. We have also seen that propositions are expressions of type *Prop*. In the formal system, if *P* is a proposition, a proof of *P* is just an expression of type *P*. This is the final piece of the puzzle: we use Lean as a proof assistant by writing down a proposition *P*, writing down an expression *p*, and asking Lean to confirm that *p* has type *P*. The fact that *2 + 2 = 4* has an easy proof, that we will explain later:

```
theorem two_plus_two_is_four : 2 + 2 = 4 := rfl
```

In contrast, proving Fermat's last theorem is considerably harder.

```
theorem Fermat_last_theorem : Fermat_statement := sorry
```

Lean knows that *sorry* is not a real proof, and it flags a warning there. If you manage to replace *sorry* by a real Lean expression, please let us know. We will be very impressed.

So, in Lean, one can write programs and execute them, and one can state propositions and prove them. In fact, one can state propositions about programs and then prove those statements as well. This is known as *software verification*; it is a means of obtaining a strong guarantee that a computer program behaves as intended, something that is important, say, if you are using the software to control a nuclear reactor or fly an airplane.

This course is not about software verification. We will be using Lean 4 primarily as a programming language, one in which we can easily define logical expressions and manipulate them. To a small extent, we will also write some simple proofs in Lean. This will help us think about proof systems and rules, and understand how they work. Taken together, these two activities embody the general vision that animates this course: knowing how to work with formally specified expressions and rules opens up a world of opportunity. It is the key to unlocking the secrets of the universe.

3.2 Using Lean as a functional programming language

There is a preliminary user's manual for Lean, still a work in progress, here. The fact that Lean is a functional programming language means that instead of presenting a program as a list of instructions, you simply *define* functions and ask Lean to evaluate them.

```
def foo n := 3 * n + 7

#eval foo 3
#eval foo (foo 3)

def bar n := foo (foo n) + 3

#eval bar 3
#eval bar (bar 3)
```

There is no global state: any value a function can act on is passed as an explicit argument and is never changed. For that reason, functional programming languages are amenable to parallelization.

Nonetheless, Lean can do handle system IO using the *IO monad*, and can accommodate an imperative style of programming using *do notation*.

```
def printExample : IO Unit:= do
  IO.println "hello"
  IO.println "world"

#eval printExample
```

Recursive definitions are built into Lean.

```
def factorial : Nat → Nat
  | 0       => 1
  | (n + 1) => (n + 1) * factorial n

#eval factorial 10
#eval factorial 100
```

Here is a solution to the Towers of Hanoi problem:

```
def hanoi (numPegs start finish aux : Nat) : IO Unit :=
  match numPegs with
  | 0     => pure ()
  | n + 1 => do
      hanoi n start aux finish
      IO.println s!"Move disk {n + 1} from peg {start} to peg {finish}"
      hanoi n aux finish start

#eval hanoi 7 1 2 3
```

You can also define things by recursion on lists:

```
def addNums : List Nat → Nat
  | []     => 0
  | a::as => a + addNums as

#eval addNums [0, 1, 2, 3, 4, 5, 6]
```

In fact, there are a number of useful functions built into Lean's library. The function *List.range n* returns the list *[0, 1, …, n-1]*, and the functions *List.map* and *List.foldl* and *List.foldr* implement the usual map and fold functions for lists. By opening the *List* namespace, we can refer to these as *range*, *map*, *foldl*, and *foldr*. In the examples below, the dollar sign has the same effect as putting parentheses around everything that appears afterward.

```
#eval List.range 7

section
open List

#eval range 7
#eval addNums $ range 7
#eval map (fun x => x + 3) $ range 7
#eval foldl (. + .) 0 $ range 7

end
```

The scope of the *open* command is limited to the section, and the cryptic inscription *(. + .)* is notation for the addition function. Lean also supports projection notation that is useful when the corresponding namespace is not open:

```
def myRange := List.range 7
#eval myRange.map fun x => x + 3
```

Because *myRange* has type *List Nat*, Lean interprets *myrange.map fun x => x + 3* as *List.map (fun x => x + 3) myrange*. In other words, it automatically interprets *map* as being in the *List* namespace, and then it interprets *myrange* as the first *List* argument.

This course assumes you have some familiarity with functional programming. There is a free online textbook, Learn You a Haskell for Great Good that you might find helpful; porting some of the examples there to Lean is a good exercise. We will all suffer from the fact that documentation for Lean 4 barely exists at the moment, but we will do our best to provide you with enough examples for you to be able to figure out how to do what you need to do. One trick is to nose around the Lean code base itself. If you ctrl-click on the name of a function in the Lean library, the editor will jump to the definition, and you can look around and see what else is there. Another strategy is simply to ask us, ask each other, or ask questions on the Lean Zulip chat. We are all in this together.

When working with a functional programming language, there are often clever tricks for doing things that you may be more comfortable doing in an imperative programming language. For example, as explained in Section 2.3, here are Lean's definitions of the *reverse* and *append* functions for lists:

```
namespace hidden

def reverseAux : List α → List α → List α
  | [],    r => r
  | a::l, r => reverseAux l (a::r)

def reverse (as : List α) :List α :=
  reverseAux as []

protected def append (as bs : List α) : List α :=
  reverseAux as.reverse bs

end hidden
```

The function *reverseAux l r* reverses the elements of list *l* and adds them to the front of *r*. When called from *reverse l*, the argument *r* acts as an *accumulator*, storing the partial result. Because *reverseAux* is tail recursive, Lean's compiler can implement it efficiently as a loop rather than a recursive function. We have defined these functions in a namespace named *hidden* so that they don't conflict with the ones in Lean's library if you open the *List* namespace.

In Lean's foundation, every function is totally defined. In particular, every function that Lean computes has to terminates (in principle) on every input. Lean 4 will eventually support arbitrary recursive definitions in which the arguments in a recursive call decrease by some measure, but some work is needed to justify these calls in the underlying foundation. In the meanwhile, we can always cheat by using the *partial* keyword, which will let us perform arbitrary recursive calls.

```
partial def gcd m n :=
  if n = 0 then m else gcd n (m % n)

#eval gcd 45 30
#eval gcd 37252 49824
```

Using *partial* takes us outside the formal foundation; Lean will not let us prove anything about *gcd* when we define it this way. Using *partial* also makes it easy for us to shoot ourselves in the foot:

```
partial def bad (n : Nat) : Nat := bad (n + 1)
```

On homework exercises, you should try to use structural recursion when you can, but don't hesitate to use *partial* whenever Lean complains about a recursive definition. We will not penalize you for it.

The following definition of the Fibonacci numbers does not require the *partial* keyword:

```
def fib' : Nat → Nat
  | 0 => 0
  | 1 => 1
  | n + 2 => fib' (n + 1) + fib' n
```

But it is inefficient; you should convince yourself that the natural evaluation strategy requires exponential time. The following definition avoids that.

```
def fibAux : Nat → Nat × Nat
  | 0     => (0, 1)
  | n + 1 => let p := fibAux n
             (p.2, p.1 + p.2)

def fib n := (fibAux n).1

#eval (List.range 20).map fib
```

Producing a *list* of Fibonacci numbers, however, as we have done here is inefficient; you should convince yourself that the running time is quadratic. In the exercises, we ask you to define a function that computes a list of Fibonacci values with running time linear in the length of the list.

3.3 Inductive data types in Lean

One reason that computer scientists and logicians tend to like functional programming languages is that they often provide good support for defining inductive data types and then using structural recursion on such types. For example, here is a Lean definition of the extended binary trees that we defined in mathematical terms in Section 2.3:

```
import Init

inductive BinTree
  | empty : BinTree
  | node  : BinTree → BinTree → BinTree
  deriving Repr, DecidableEq, Inhabited

open BinTree
```

The command *import Init* imports a part of the initial library for us to use. The command *open BinTree* allows us to write *empty* and *node* instead of *BinTree.empty* and *BinTree.node*. Note the Lean convention of capitalizing the names of data types.

The last line of the definition, the one that begins with the word *deriving*, is boilerplate. It tells Lean to automatically generate a few additional functions that are useful. The directive *deriving Repr* tells Lean to define an internal function that can be used to represent any *BinTree* as a string. This is the string that is printed out by any *#eval* command whose argument evaluates to a *BinTree*. Adding *DecidableEq* defines a function that tests whether two *BinTrees* are equal, and adding *Inhabited* defines an arbitrary value of the data type to serve as a default value for function that need one. The following illustrates their use.

```
#eval node empty (node empty empty)

#eval empty == node empty empty   -- evaluates to false

#eval (arbitrary : BinTree)   -- BinTree.empty
```

We can now define the functions *size* and *depth* by structural recursion:

```
def size : BinTree → Nat
  | empty    => 0
  | node a b => 1 + size a + size b

def depth : BinTree → Nat
  | empty    => 0
  | node a b => 1 + Nat.max (depth a) (depth b)

def example_tree := node (node empty empty) (node empty (node empty empty))

#eval size example_tree
#eval depth example_tree
```

In fact, the *List* data type is also inductively defined.

```
#print List
```

You should try writing the inductive definition on your own. Call it *MyList*, and then try *#print MyList* to see how it compares.

3.4 Using Lean as an imperative programming language

The fact that Lean is a functional programming language means that there is no global notion of *state*. Functions take values as input and return values as output; there are no global or even local variables that are changed by the result of a function call.

But one of the interesting features of Lean is a functional programming language is that it incorporates features that make it *feel* like an imperative programming language. The following example shows how to print out, for each value i less than 100, the the sum of the numbers up to i.

```
def showSums : IO Unit := do
  let mut sum := 0
  for i in [0:100] do
    sum := sum + i
    IO.println s!"i: {i}, sum: {sum}"

#eval showSums
```

You can use a loop not just to print values, but also to compute values. The following is a boolean test for primality:

```
def isPrime (n : Nat) : Bool := do
  if n < 2 then false else
    for i in [2:n] do
      if n % i = 0 then
        return false
      if i * i > n then
        return true
    true
```

You can use such a function with the list primitives to construct a list of the first 10,000 prime numbers.

```
#eval (List.range 10000).filter isPrime
```

But you can also use it with Lean's support for *arrays*. Within the formal foundation these are modeled as lists, but the compiler implements them as dynamic arrays, and for efficiency it will modify values rather than copy them whenever the old value is not referred to by another part of an expression.

```
def primes (n : Nat) : Array Nat := do
  let mut result := #[]
  for i in [2:n] do
    if isPrime i then
      result := result.push i
  result

#eval (primes 10000).size
```

Notice the notation: *#[]* denotes a fresh array (Lean infers the type from context), and the *Array.push* function adds a new element at the end of the array.

The following example shows how to compute a two-dimensional array and print it out.

```
def mulTable (n : Nat) : Array (Array Nat) := do
  let mut table := #[]
  for i in [:n] do
    let mut row := #[]
    for j in [:n] do
      row := row.push ((i + 1) * (j + 1))
    table := table.push row
  table

#eval mulTable 10

def printMulTable (n : Nat) : IO Unit := do
  let t := mulTable n
  for i in [:n] do
    for j in [:n] do
      IO.print s!"{t[i][j]} "
    IO.println ""

#eval printMulTable 10
```

3.5 Exercises

1. Using operations on *List*, write a Lean function that for every n returns the list of all the divisors of n that are less than n.

2. A natural number n is *perfect* if it is equal to the sum of the divisors less than n. Write a Lean function (with return type *Bool*) that determines whether a number n is perfect. Use it to find all the perfect numbers less than 1,000.

3. Define a recursive function *sublists*(ℓ) that, for every list ℓ, returns a list of all the sublists of ℓ. For example, given the list $[1, 2, 3]$, it should compute the list

$$[[], [1], [2], [3], [1, 2], [1, 3], [2, 3], [1, 2, 3]].$$

The elements need not be listed in that same order.

4. Prove in Lean that the length of *sublists*(ℓ) is $2^{length(\ell)}$.

5. Define a function *permutations*(ℓ) that returns a list of all the permutations of ℓ.

6. Prove in Lean that the length of *permutations*(ℓ) is *factorial*(*length*(ℓ)).

7. Define in Lean a function that, assuming ℓ is a list of lists representing an $n \times n$ array, returns a list of lists representing the transpose of that array.

8. Write a program that solves the Tower of Hanoi problem with n disks on the assumption that disks can only be moved to an *adjacent* peg. (See Section 2.5.)

9. Write a program that solves the Tower of Hanoi problem with n disks on the assumption that disks can only be moved clockwise. (See Section 2.5.)

10. Define a Lean data type of binary trees in which every node is numbered by a label. Define a Lean function to compute the sum of the nodes in such a tree. Also write functions to list the elements in a preorder, postorder, and inorder traversal.

PROPOSITIONAL LOGIC

We are finally ready to turn to the proper subject matter of this course, logic. We will see that although propositional logic has limited expressive power, it can be used to carry out useful combinatorial reasoning in a wide range of applications.

4.1 Syntax

We start with a stock of variables $p_0, p_1, p_2, \ldots$ that we take to range over propositions, like "the sky is blue" or "2 + 2 = 5". We'll make the interpretation of propositional logic explicit in the next section, but, intuitively, propositions are things that can be either true or false. (More precisely, this is the *classical* interpretation of propositional logic, which is the one we will focus on in this course.) Each propositional variable is a *formula*, and we also include symbols $\top$ and $\bot$ for "true" and "false" respectively. We also provide means for building new formulas from old ones. The following is a paradigm instance of an inductive definition.

Definition

The set of propositional formulas is generated inductively as follows:

- Each variable p_i is a formula.

- $\top$ and $\bot$ are formulas.

- If A is a formula, so is $\neg A$ ("not A").

- If A and B are formulas, so are

    - $A \wedge B$ ("A and B"),

    - $A \vee B$ ("A or B"),

    - $A \rightarrow B$ ("A implies B"), and

    - $A \leftrightarrow B$ ("A if and only if B").

We will see later that there is some redundancy here; we could get by with fewer connectives and define the others in terms of those. Conversely, there are other connectives that can be defined in terms of these. But the ones we have included form a *complete* set of connectives, which is to say, any conceivable connective can be defined in terms of these, in a sense we will clarify later.

The fact that the set is generated inductively means that we can use recursion to define functions on the set of propositional

formulas, as follows:

$$complexity(p_i) = 0$$
$$complexity(\top) = 0$$
$$complexity(\bot) = 0$$
$$complexity(\neg A) = complexity(A) + 1$$
$$complexity(A \wedge B) = complexity(A) + complexity(B) + 1$$
$$complexity(A \vee B) = complexity(A) + complexity(B) + 1$$
$$complexity(A \rightarrow B) = complexity(A) + complexity(B) + 1$$
$$complexity(A \leftrightarrow B) = complexity(A) + complexity(B) + 1$$

The function $complexity(A)$ counts the number of connectives. The function $depth(A)$, defined in a similar way, computes the depth of the parse tree.

$$depth(p_i) = 0$$
$$depth(\top) = 0$$
$$depth(\bot) = 0$$
$$depth(\neg A) = depth(A) + 1$$
$$depth(A \wedge B) = \max(depth(A), depth(B)) + 1$$
$$depth(A \vee B) = \max(depth(A), depth(B)) + 1$$
$$depth(A \rightarrow B) = \max(depth(A), depth(B)) + 1$$
$$depth(A \leftrightarrow B) = \max(depth(A), depth(B)) + 1$$

Here's an example of a proof by induction:

Theorem

For every formula A, we have $complexity(A) \leq 2^{depth(A)} - 1$.

Proof

In the base case, we have

$$complexity(p_i) = 0 = 2^0 - 1 = 2^{depth(p_i)} - 1,$$

and similarly for $\top$ and $\bot$. In the case for negation, assuming the claim holds of A, we have

$$complexity(\neg A) = complexity(A) + 1$$
$$\leq 2^{depth(A)} - 1 + 1$$
$$\leq 2^{depth(A)} + 2^{depth(A)} - 1$$
$$\leq 2^{depth(A)+1} - 1$$
$$= 2^{depth(\neg A)} - 1.$$

Finally, assuming the claim holds of A and B, we have

$$complexity(A \wedge B) = complexity(A) + complexity(B) + 1$$
$$\leq 2^{depth(A)} - 1 + 2^{depth(B)} - 1 + 1$$
$$\leq 2 \cdot 2^{\max(depth(A),depth(B))} - 1$$
$$= 2^{\max(depth(A),depth(B))+1} - 1$$
$$= 2^{depth(A \wedge B)} - 1,$$

---

and similarly for the other binary connectives.

In our metatheory, we will use variables $p, q, r, \ldots$ to range over propositional variables, and $A, B, C, \ldots$ to range over propositional formulas. The formulas p_i, $\top$, and $\bot$ are called *atomic* formulas. If A is a formula, B is a *subformula* of A if B occurs somewhere in A. We can make this precise by defining the set of subformulas of any formula A inductively as follows:

$$
\begin{aligned}
subformulas(A) &= \{A\} \quad \text{if } A \text{ is atomic} \\
subformulas(\neg A) &= \{\neg A\} \cup subformulas(A) \\
subformulas(A \star B) &= \{A \star B\} \cup subformulas(A) \cup subformulas(B)
\end{aligned}
$$

In the last clause, the star is supposed to represent any binary connective.

If A and B are formulas and p is a propositional variable, the notation $A[B/p]$ denotes the result of substituting B for p in A. Beware: the notation for this varies widely; $A[p \mapsto B]$ is also becoming common in computer science. The meaning is once again given by a recursive definition:

$$
\begin{aligned}
p_i[B/p] &= \begin{cases} B & \text{if } p \text{ is } p_i \\ p_i & \text{otherwise} \end{cases} \\
(\neg C)[B/p] &= \neg(C[B/p]) \\
(C \star D)[B/p] &= C[B/p] \star D[B/p]
\end{aligned}
$$

4.2 Semantics

Consider the formula $p \wedge (\neg q \vee r)$. Is it true? Well, that depends on the propositions p, q, and r. More precisely, it depends on whether they are true — and, in fact, that is all it depends on. In other words, once we specify which of p, q, and r are true and which are false, the truth value of $p \wedge (\neg q \vee r)$ is completely determined.

To make this last claim precise, we will use the set $\{\top, \bot\}$ to represent the truth values *true* and *false*. It doesn't really matter what sorts of mathematical objects those are, as long as they are distinct. You can take them to be the corresponding propositional formulas, or you can take $\top$ to be the number 1 and $\bot$ to be the number 0. A *truth assignment* is a function from propositional variables to the set $\{\top, \bot\}$, that is, a function which assigns a value of true or false to each propositional variable. Any truth assignment v extends to a function $\bar{v}$ that assigns a value of $\top$ or $\bot$ to any propositional formula. It is defined recursively as follows:

$$
\begin{aligned}
\bar{v}(p_i) &= v(p_i) \\
\bar{v}(\top) &= \top \\
\bar{v}(\bot) &= \bot \\
\bar{v}(\neg A) &= \begin{cases} \top & \text{if } \bar{v}(A) = \bot \\ \bot & \text{otherwise} \end{cases} \\
\bar{v}(A \wedge B) &= \begin{cases} \top & \text{if } \bar{v}(A) = \top \text{ and } \bar{v}(B) = \top \\ \bot & \text{otherwise} \end{cases} \\
\bar{v}(A \vee B) &= \begin{cases} \top & \text{if } \bar{v}(A) = \top \text{ or } \bar{v}(B) = \top \\ \bot & \text{otherwise} \end{cases} \\
\bar{v}(A \to B) &= \begin{cases} \top & \text{if } \bar{v}(A) = \bot \text{ or } \bar{v}(B) = \top \\ \bot & \text{otherwise} \end{cases} \\
\bar{v}(A \leftrightarrow B) &= \begin{cases} \top & \text{if } \bar{v}(A) = \bar{v}(B) \\ \bot & \text{otherwise} \end{cases}
\end{aligned}
$$

It is common to write $[\![A]\!]_v$ instead of $\bar{v}(A)$. Double-square brackets like these are often used to denote a semantic value that is assigned to a syntactic object. Think of $[\![A]\!]_v$ as giving the *meaning* of A with respect to the interpretation given by v. In this case, variables are interpreted as standing for truth values and the meaning of the formula is the resulting truth value, but in the chapters to come we will come across other semantic interpretations of this sort.

The following definitions are now fundamental to logic. Make sure you are clear on the terminology and know how to use it. If you can use these terms correctly, you can pass as a logician. If you get the terminology wrong, you'll be frowned upon.

- If $[\![A]\!]_v = \top$, we say that A is *satisfied* by v, or that v is a *satisfying assignment* for A. We also sometimes write $\models_v A$.

- A formula A is *satisfiable* if there is some truth assignment that satisfies it. A formula A is *unsatisfiable* if it is not satisfiable.

- A formula A is *valid*, or a *tautology* if it is satisfied by *every* truth assignment. In other words, A is valid if $[\![A]\!]_v = \top$ for every truth assignment v.

- If Γ is a set of propositional formulas, we say that Γ is *satisfied by* v if every formula in Γ is satisfied by v. In other words, Γ is satisfied by v if $[\![A]\!]_v = \top$ for every A in Γ.

- A set of formulas Γ is *satisfiable* if it is satisfied by some truth assignment v. Otherwise, it is *unsatisfiable*.

- If Γ is a set of propositional formulas and A is a propositional formula, we say Γ entails A if every truth assignment that satisfies Γ also satisfies A. Roughly speaking, this says that whenever the formulas in Γ are true, then A is also true. In this case, A is also said to be a *logical consequence* of Γ.

- Two formulas A and B are *logically equivalent* if each one entails the other, that is, we have $\{A\} \models B$ and $\{B\} \models A$. When that happens, we write $A \equiv B$.

There is a lot to digest here, but it is important that you become comfortable with these definitions. The mathematical analysis of truth and logical consequence is one of the crowning achievements of modern logic, and this basic framework for reasoning about expressions and their meaning has been applied to countless other settings in logic and computer science.

You should also get used to using semantic notions in proofs. For example:

Theorem

A propositional formula A is valid if and only if $\neg A$ is unsatisfiable.

Proof

A is valid if and only if $[\![A]\!]_v = \top$ for every truth assignment v. By the definition of $[\![\neg A]\!]_v$, this happens if and only if $[\![\neg A]\!]_v = \bot$ for every v, which is the same as saying that $\neg A$ is unsatisfiable.

4.3 Calculating with propositions

Remember that Leibniz imagined that one day we would be able to calculate with propositions. What he had noticed is that propositions, like numbers, obey algebraic laws. Here are some of them:

- $A \vee \neg A \equiv \top$
- $A \wedge \neg A \equiv \bot$
- $\neg \neg A \equiv A$

- $A \vee A \equiv A$

- $A \wedge A \equiv A$

- $A \vee \bot \equiv A$

- $A \wedge \bot \equiv \bot$

- $A \vee \top \equiv \top$

- $A \wedge \top \equiv A$

- $A \vee B \equiv B \vee A$

- $A \wedge B \equiv B \wedge A$

- $(A \vee B) \vee C \equiv A \vee (B \vee C)$

- $(A \wedge B) \wedge C \equiv A \wedge (B \wedge C)$

- $\neg(A \wedge B) \equiv \neg A \vee \neg B$

- $\neg(A \vee B) \equiv \neg A \wedge \neg B$

- $A \wedge (B \vee C) \equiv (A \wedge B) \vee (A \wedge C)$

- $A \vee (B \wedge C) \equiv (A \vee B) \wedge (A \vee C)$

- $A \wedge (A \vee B) \equiv A$

- $A \vee (A \wedge B) \equiv A$

The equivalences $\neg(A \wedge B) \equiv \neg A \vee \neg B$ and $\neg(A \vee B) \equiv \neg A \wedge \neg B$ are known as *De Morgan's laws*. It is not hard to show that all the logical connectives respect equivalence, and hence substituting equivalent formulas for a variable in a formula results in equivalent formulas. This means that, as Leibniz imagined, we can prove that a Boolean formula is valid by calculating to show that it is equivalent to $\top$. Here is an example.

Theorem

For any propositional formulas A and B, we have $(A \wedge \neg B) \vee B \equiv A \vee B$.

Proof

$$(A \wedge \neg B) \vee B \equiv (A \vee B) \wedge (\neg B \vee B)$$
$$\equiv (A \vee B) \wedge \top$$
$$\equiv (A \vee B).$$

Mathematicians have a trick, called *quotienting*, for turning an equivalence relation into an equality. If we interpret A, B, and C as *equivalence classes* of formulas instead of formulas, the equivalences listed above become identities. The resulting algebraic structure is known as a *Boolean algebra*, and we can view the preceding proof as establishing an identity that holds in any Boolean algebra. The same trick is used, for example, to interpret an equivalence between numbers modulo 12, like $5 + 9 \equiv 2$ as an identity on the structure $\mathbb{Z}/12\mathbb{Z}$.

4.4 Complete sets of connectives

You may have noticed that our choice of connectives is redundant. For example, the following equivalences show that we can get by with $\neg$, $\vee$, and $\bot$ alone:

$$A \wedge B \equiv \neg(\neg A \vee \neg B)$$
$$A \rightarrow B \equiv \neg A \vee B$$
$$A \leftrightarrow B \equiv (A \rightarrow B) \wedge (B \rightarrow A)$$
$$\top \equiv \neg\bot$$

We can even define $\bot$ as $P \wedge \neg P$ for any propositional variable P, though that has the sometimes annoying consequence that we cannot express the constants $\top$ and $\bot$ without using a propositional variable.

Let $f(x_0, \ldots, x_{n-1})$ be a function that takes n truth values and returns a truth value. A formula A with variables $p_0, \ldots, p_{n-1}$ is said to *represent* f if for every truth assignment v,

$$[\![A]\!]_v = f(v(p_0), \ldots, v(p_{n-1})).$$

If you think of f as a truth table, this says that the truth table of A is f.

A set of connectives is said to be *complete* if every function f is represented by some formula A involving those connectives. In class, we'll discuss how to prove that $\{\wedge, \neg\}$ is a complete set of connectives in that sense.

It is now straightforward to show that a certain set of connectives is complete: just show how to define $\vee$ and $\neg$ in terms of them. Showing that a set of connectives is *not* complete typically requires some more ingenuity. One idea, as suggested in Section 2.4, is to look for some invariant property of the formulas that *are* represented.

4.5 Normal forms

For both theoretical reasons and practical reasons, it is often useful to know that formulas can be expressed in particularly simple or convenient forms.

Definition

An *atomic* formula is a variable or one of the constants $\top$ or $\bot$. A *literal* is an atomic formula or a negated atomic formula.

Definition

The set of propositional formulas in *negation normal form* (NNF) is generated inductively as follows:

- Each literal is in negation normal form.

- If A and B are in negation normal form, then so are $A \wedge B$ and $A \vee B$.

More concisely, the set of formulas in negation normal form is the smallest set of formulas containing the literals and closed under conjunction and disjunction. If we identify $\top$ with $\neg\bot$ and $\neg\top$ with $\bot$, we can alternatively characterize the formulas in negation normal form as the smallest set of formulas containing $\top$, $\bot$, variables, and their negations, and closed under conjunction and disjunction.

Proposition

Every propositional formula is equivalent to one in negation normal form.

Proof

First use the identities $A \leftrightarrow B \equiv (A \rightarrow B) \wedge (B \rightarrow A)$ and $A \rightarrow B \equiv \neg A \vee B$ to get rid of $\leftrightarrow$ and $\rightarrow$. Then use De Morgan's laws together with $\neg\neg A \equiv A$ to push negations down to the atomic formulas.

More formally, we can prove by induction on propositional formulas A that both A and $\neg A$ are equivalent to formulas in negation normal form. (You should try to write that proof down carefully.) Putting a formula in negation normal form is reasonably efficient. You should convince yourself that if A is in negation normal form, then putting $\neg A$ in negation normal form amounts to switching all the following in A:

- $\top$ with $\bot$

- variables p_i with their negations $\neg p_i$

- $\wedge$ with $\vee$.

We will see that *conjunctive normal form* (CNF) and *disjunctive normal form* (DNF) are also important representations of propositional formulas. A formula is in conjunctive normal form if it can be written as conjunction of disjunctions of literals, in other words, if it can be written as a big "and" of "or" expressions:

$$\bigwedge_{i<n} \left(\bigvee_{j<m_i} \pm \ell_j \right).$$

where each ℓ_j is a literal. Here is an example:

$$(p \vee \neg q \vee r) \wedge (\neg p \vee s) \wedge (\neg r \vee s \vee \neg t).$$

We can think of $\bot$ as the empty disjunction (because a disjunction is true only when one of the disjuncts is true) and we can think of $\top$ as the empty conjunction (because a conjunction is true when all of its conjuncts are true, which happens trivially when there aren't any).

Dually, a formula is in disjunctive normal form if it is an "or" and "and" expressions:

$$\bigvee_{i<n} \left(\bigwedge_{j<m_i} \pm \ell_j \right).$$

If you switch $\wedge$ and $\vee$ in the previous example, you have a formula in disjunctive normal form.

It is pretty clear that if you take the conjunction of two formulas in CNF the result is a CNF formula (modulo associating parentheses), and, similarly, the disjunction of two formulas in DNF is again DNF. The following is less obvious:

Lemma

The disjunction of two CNF formulas is equivalent to a CNF formula, and dually for DNF formulas.

Proof

For the first claim, we use the equivalence $A \vee (B \wedge C) \equiv (A \vee B) \wedge (A \vee C)$. By induction on n, we have that for every sequence of formulas $B_0, \ldots, B_{n-1}$ we have $A \vee \bigwedge_{i<n} B_i \equiv \bigwedge_{i<n} (A \vee B_i)$. Then by induction on n' we have $\bigwedge_{i'<n'} A_{i'} \vee \bigwedge_{i<n} B_i \equiv \bigwedge_{i'<n'} \bigwedge_{i<n} (A_{i'} \vee B_i)$. Since each $A_{i'}$ and each B_i is a disjunction of literals, this yields the result.

The second claim is proved similarly, switching $\wedge$ and $\vee$.

Proposition

Every propositional formula is equivalent to one in conjunctive normal form, and also to one in disjunctive normal form.

Proof

Since we already know that every formula is equivalent to one in negation normal form, we can use induction on that set of formulas. The claim is clearly true of $\top$, $\bot$, p_i, and $\neg p_i$. By the previous lemma, whenever it is true of A and B, it is also true of $A \wedge B$ and $A \vee B$.

In contrast to putting formulas in negation normal form, the exercises below show that the smallest CNF or DNF equivalent of a formula A can be exponentially longer than A.

We will see that conjunctive normal form is commonly used in automated reasoning. Notice that if a disjunction of literals contains a duplicated literal, deleting the duplicate results in an equivalent formula. We can similarly delete any occurrence of $\bot$. A disjunction of literals is called a *clause*. Since the order of the disjuncts and repetitions don't matter, we generally identify clauses with the corresponding set of literals; for example, the clause $p \vee \neg q \vee r$ is associated with the set $\{p, \neg q, r\}$. If a clause contains a pair p_i and $\neg p_i$, or if it contains $\top$, it is equivalent to $\top$. If Γ is a set of clauses, we think of Γ as saying that all the clauses in Γ are true. With this identification, every formula in conjunctive normal form is equivalent to a set of clauses. An empty clause corresponds to $\bot$, and an empty set of clauses corresponds to $\top$.

The dual notion to a clause is a conjunction like $\neg p \wedge q \wedge \neg r$. If each variable occurs at most once (either positively or negatively), we can think of this as a *partial truth assignment*. In this example, any truth assignment that satisfies the formula has to set p false, q true, and r false.

4.6 Exercises

1. Prove that if A is a subformula of B and B is a subformula of C then A is a subformula of C. (Hint: prove by induction on C that for every $B \in subformulas(C)$, every subformula of B is a subformula of C.)

2. Prove that for every A, B, and p, $depth(A[B/p]) \leq depth(A) + depth(B)$.

3. Prove that A and B are logically equivalent if and only if the formula $A \leftrightarrow B$ is valid.

4. Use algebraic calculations to show that all of the following are tautologies:

   - $((A \wedge \neg B) \vee B) \leftrightarrow (A \vee B)$

   - $(A \rightarrow \neg A) \rightarrow \neg A$

   - $(A \rightarrow B) \leftrightarrow (\neg B \rightarrow \neg A)$

   - $A \rightarrow (B \rightarrow A \wedge B)$

5. The *Sheffer stroke* $A \mid B$, also known as "nand," says that A and B are not both true. Show that $\{\mid\}$ is a complete set of connectives. Do the same for "nor," that is, the binary connective that holds if neither A nor B is true.

6. Show that $\{\wedge, \neg\}$ and $\{\rightarrow, \bot\}$ are complete sets of connectives.

7. Show that $\{\rightarrow, \vee, \wedge\}$ is not a complete set of connectives. Conclude that $\{\rightarrow, \vee, \wedge, \leftrightarrow, \top\}$ is not a complete set of connectives.

8. Show that $\{\bot, \leftrightarrow\}$ is not a complete set of connectives. Conclude that $\{\bot, \top, \neg, \leftrightarrow, \oplus\}$ is not complete. Here $A \oplus B$ is the "exclusive or," which is to say, $A \oplus B$ is true if one of A or B is true but not both.

9. Using the property $A \vee (B \wedge C) \equiv (A \vee B) \wedge (A \vee C)$ and the dual statement with $\wedge$ and $\vee$ switched, put $(p_1 \wedge p_2) \vee (q_1 \wedge q_2) \vee (r_1 \wedge r_2)$ in conjunctive normal form.

---

10. The boolean function $parity(x_0, x_1, \ldots, x_{n-1})$ holds if and only if an odd number of the x_i s are true. It is represented by the formula $p_0 \oplus p_1 \oplus \cdots \oplus p_{n-1}$. Show that any CNF formula representing the parity function has to have at least 2^n clauses.

IMPLEMENTING PROPOSITIONAL LOGIC

5.1 Propositional formulas

We have seen that the set of propositional formulas can be defined inductively, and we have seen that Lean makes it easy to specify inductively defined types. It's a match made in heaven! Here is the definition of the type of propositional formulas that we will use in this course:

```
namespace hidden

inductive PropForm
  | tr     : PropForm
  | fls    : PropForm
  | var    : String → PropForm
  | conj   : PropForm → PropForm → PropForm
  | disj   : PropForm → PropForm → PropForm
  | impl   : PropForm → PropForm → PropForm
  | neg    : PropForm → PropForm
  | biImpl : PropForm → PropForm → PropForm
  deriving Repr, DecidableEq

end hidden

#print PropForm

open PropForm

#check (impl (conj (var "p") (var "q")) (var "r"))
```

You can find this example in the file *implementing_propositional_logic/examples.lean* in the *User* folder of the course repository.

The command *import LAMR.Util.Propositional* at the top of the file imports the part of the library with functions that we provide for you to deal with propositional logic. We will often put a copy of a definition from the library in an examples file so you can experiment with it. Here we have put it in a namespace called *hidden* so that our copy's full name is *hidden.PropForm*, which won't conflict with the one in the library. Outside the *hidden* namespace, the command *#print PropForm* refers to the real one, that is, the one in the library. The command *open PropForm* means that we can write, for example, *tr* for the first constructor instead of *PropForm.tr*. Try writing some propositional formulas of your own. There should be squiggly blue lines under the the *#print* and *#check* commands in VSCode, indicating that there is Lean output associated with these. You can see it by hovering over the commands, or by moving the caret to the command and checking the *Lean infoview* window.

The phrase *deriving Repr, DecidableEq* tells Lean to automatically define functions to be used to test equality of two expressions of type *PropForm* and to display the result of an *#eval*. We'll generally leave these out of the display from now on. You an always use *#check* and *#print* to learn more about a definition in the library. If you hold down *ctrl* and

click on an identifier, the VSCode Lean extension will take you to the definition in the library. Simply holding down *ctrl* and hovering over it will show you the definition in a pop-up window. Try taking a look at the definition of *PropForm* in the library.

Writing propositional formulas using constructors can be a pain in the neck. In the library, we have used Lean's mechanisms for defining new syntax to implement nicer syntax.

```
#check prop!{p ∧ q → (r ∨ ¬ p) → q}
#check prop!{p ∧ q ∧ r → p}

def propExample := prop!{p ∧ q → r ∧ p ∨ ¬ s1 → s2 }

#print propExample
#eval propExample

#eval toString propExample
```

You can get the symbols by typing \and, \to, \or, \not, and \iff in VS Code. And, in general, when you see a symbol in VSCode, hovering over it with the mouse shows you how to type it. Once again, try typing some examples of your own. The library defines the function *PropForm.toString* that produces a more readable version of a propositional formula, one that, when inserted within the *prop!{…}* brackets, should produce the same result.

Because *PropForm* is inductively defined, we can easily define functions using structural recursion.

```
namespace PropForm

def complexity : PropForm → Nat
  | var _ => 0
  | tr => 0
  | fls => 0
  | neg A => complexity A + 1
  | conj A B => complexity A + complexity B + 1
  | disj A B => complexity A + complexity B + 1
  | impl A B => complexity A + complexity B + 1
  | biImpl A B => complexity A + complexity B + 1

def depth : PropForm → Nat
  | var _ => 0
  | tr => 0
  | fls => 0
  | neg A => depth A + 1
  | conj A B => Nat.max (depth A) (depth B) + 1
  | disj A B => Nat.max (depth A) (depth B) + 1
  | impl A B => Nat.max (depth A) (depth B) + 1
  | biImpl A B => Nat.max (depth A) (depth B) + 1

def vars : PropForm → List String
  | var s => [s]
  | tr => []
  | fls => []
  | neg A => vars A
  | conj A B => (vars A).union (vars B)
  | disj A B => (vars A).union (vars B)
  | impl A B => (vars A).union (vars B)
  | biImpl A B => (vars A).union (vars B)

#eval complexity propExample
#eval depth propExample
```

```
#eval vars propExample

end PropForm

#eval PropForm.complexity propExample
#eval propExample.complexity
```

The function *List.union* returns concatenation of the two lists with duplicates removed, assuming that the original two lists had no duplicate elements.

5.2 Semantics

The course library defines the type *PropAssignment* to be *List (String × Bool)*. If *v* has type *PropAssignment*, you should think of the expression *v.eval s* as assigning a truth value to the variable named *s*. The following function then evaluates the truth value of any propositional formula under assignment *v*:

```
def PropForm.eval (v : PropAssignment) : PropForm → Bool
  | var s => v.eval s
  | tr => true
  | fls => false
  | neg A => !(eval v A)
  | conj A B => (eval v A) && (eval v B)
  | disj A B => (eval v A) || (eval v B)
  | impl A B => !(eval v A) || (eval v B)
  | biImpl A B => (!(eval v A) || (eval v B)) && (!(eval v B) || (eval v A))

-- try it out
#eval let v := PropAssignment.mk [("p", true), ("q", true), ("r", true)]
      propExample.eval v
```

The example at the end defines *v* to be the assignment that assigns the value *true* to the strings *"p"*, *"q"*, and *"r"* and false to all the others. This is a reasonably convenient way to describe truth assignments manually, so the library provides a function *PropAssignment.mk* and notation *propassign!{…}* to support that.

```
#check propassign!{p, q, r}

#eval propExample.eval propassign!{p, q, r}
```

You should think about how the next function manages to compute a list of all the sublists of a given list. It is analogous to the power set operation in set theory.

```
def allSublists : List α → List (List α)
  | [] => [[]]
  | (a :: as) =>
      let recval := allSublists as
      recval.map (a :: .) ++ recval

#eval allSublists propExample.vars
```

With that in hand, here is a function that computes the truth table of a propositional formula. The value of *truthTable A* is a list of pairs: the first element of the pair is the list of *true/false* values assigned to the elements of *vars A*, and the second element is the truth value of *A* under that assignment.

```
def truthTable (A : PropForm) : List (List Bool × Bool) :=
  let vars := A.vars
  let assignments := (allSublists vars).map (fun l => PropAssignment.mk (l.map (.,␣
  ↪true)))
  let evalLine := fun v : PropAssignment => (vars.map v.eval, A.eval v)
  assignments.map evalLine

#eval truthTable propExample
```

We can now use the list operation *List.all* to test whether a formula is valid, and we can use *List.some* to test whether it is satisfiable.

```
def PropForm.isValid (A : PropForm) : Bool := List.all (truthTable A) Prod.snd
def PropForm.isSat(A : PropForm) : Bool := List.any (truthTable A) Prod.snd

#eval propExample.isValid
#eval propExample.isSat
```

5.3 Normal Forms

The library defines an inductive type of negation-normal form formulas:

```
inductive Lit
  | tr  : Lit
  | fls : Lit
  | pos : String → Lit
  | neg : String → Lit

inductive NnfForm :=
  | lit  (l : Lit)         : NnfForm
  | conj (p q : NnfForm) : NnfForm
  | disj (p q : NnfForm) : NnfForm
```

It is then straightforward to define the negation operation for formulas in negation normal form, and a translation from propositional formulas to formulas in negation normal form.

```
def Lit.negate : Lit → Lit
  | tr    => fls
  | fls   => tr
  | pos s => neg s
  | neg s => pos s

def NnfForm.neg : NnfForm → NnfForm
  | lit l     => lit l.negate
  | conj p q => disj (neg p) (neg q)
  | disj p q => conj (neg p) (neg q)

namespace PropForm

def toNnfForm : PropForm → NnfForm
  | tr         => NnfForm.lit Lit.tr
  | fls        => NnfForm.lit Lit.fls
  | var n      => NnfForm.lit (Lit.pos n)
  | neg p      => p.toNnfForm.neg
```

```
  | conj p q   => NnfForm.conj p.toNnfForm q.toNnfForm
  | disj p q   => NnfForm.disj p.toNnfForm q.toNnfForm
  | impl p q   => NnfForm.disj p.toNnfForm.neg q.toNnfForm
  | biImpl p q => NnfForm.conj (NnfForm.disj p.toNnfForm.neg q.toNnfForm)
                              (NnfForm.disj q.toNnfForm.neg p.toNnfForm)

end PropForm
```

Putting the first in the namespace *NnfForm* has the effect that given *A : NnfForm*, we can write *A.neg* instead of *Nnf-Form.neg A*. Similarly, putting the second definition in the namespace *PropForm* means we can write *A.toNnfForm* to put a propositional formula in negation normal form.

We can try them out on the example defined above:

```
#eval propExample.toNnfForm
#eval toString propExample.toNnfForm
```

To handle conjunctive normal form, the library defines a type *Lit* of literals. A *Clause* is then a list of literals, and a *CnfForm* is a list of clauses.

```
def Clause := List Lit
```

```
def CnfForm := List Clause
```

As usual, you can use *#check* and *#print* to find information about them, and ctrl-click to see the definitions in the library. Since, as usual, defining things using constructors can be annoying, the library defines syntax for writing expressions of these types.

```
def exLit0 := lit!{ p }
def exLit1 := lit!{ -q }

#print exLit0
#print exLit1

def exClause0 := clause!{ p }
def exClause1 := clause!{ p -q r }
def exClause2 := clause!{ r -s }

#print exClause0
#print exClause1
#print exClause2

def exCnf0 := cnf!{
  p,
  -p q -r,
  -p q
}

def exCnf1 := cnf!{
  p -q,
  p q,
  -p -r,
  -p r
}

def exCnf2 := cnf!{
```

```
  p q,
  -p,
  -q
}

#print exCnf0
#print exCnf1
#print exCnf2

#eval toString exClause1
#eval toString exCnf2
```

Let us now consider what is needed to put an arbitrary propositional formula in conjunctive normal form. In Section 4.5, we saw that the key is to show that the disjunction of two CNF formulas is again CNF. Lean's library has a function *List.insert*, which adds an element to a list; if the element already appears in the list, it does nothing. It has a function *List.union* that will form the union of two lists; if the original two lists have no duplicates, the union won't either. Finally, we have a function *List.Union* which takes the union of a list of lists. Since clauses are lists, we can use them on clauses:

```
#eval List.insert lit!{ r } exClause0

#eval exClause0.union exClause1

#eval List.Union [exClause0, exClause1, exClause2]
```

We can now take the disjunction of a single clause and a CNF formula by taking the union of the clause with each element of the CNF formula. We can implement that with the function *List.map*:

```
#eval exCnf1.map exClause0.union
```

This applied the function "take the union with *exClause0*" to each element of *exCnf1*, and returns the resulting list. We can now define the disjunction of two CNF formulas by taking all the clauses in the first, taking the disjunction of each clause with the second CNF, and then taking the union of all of those, corresponding to the conjunctions of the CNFs. Here is the library definition, and an example:

```
def CnfForm.disj (cnf1 cnf2 : CnfForm) : CnfForm :=
(cnf1.map (fun cls => cnf2.map cls.union)).Union

#eval cnf!{p, q, u -v}.disj cnf!{r1 r2, s1 s2, t1 t2 t3}
#eval toString $ cnf!{p, q, u -v}.disj cnf!{r1 r2, s1 s2, t1 t2 t3}
```

Functional programmers like this sort of definition; it's short, clever, and inscrutable. You should think about defining the disjunction of two CNF formulas by hand, using recursions over clauses and CNF formulas. Your solution will most likely reconstruct the effect of the instance *map* and *Union* in the library definition, and that will help you understand why they make sense.

In any case, with this in hand, it is easy to define the translation from negation normal form formulas and arbitrary propositional formulas to CNF.

```
def NnfForm.toCnfForm : NnfForm → CnfForm
  | NnfForm.lit (Lit.pos s) => [ [Lit.pos s] ]
  | NnfForm.lit (Lit.neg s) => [ [Lit.neg s] ]
  | NnfForm.lit Lit.tr       => []
  | NnfForm.lit Lit.fls      => [ [] ]
  | NnfForm.conj A B         => A.toCnfForm.conj B.toCnfForm
  | NnfForm.disj A B         => A.toCnfForm.disj B.toCnfForm
```

```
def PropForm.toCnfForm (A : PropForm) : CnfForm := A.toNnfForm.toCnfForm
```

We can try them out:

```
#eval propExample.toCnfForm

#eval prop!{(p1 ∧ p2) ∨ (q1 ∧ q2)}.toCnfForm.toString

#eval prop!{(p1 ∧ p2) ∨ (q1 ∧ q2) ∨ (r1 ∧ r2) ∨ (s1 ∧ s2)}.toCnfForm.toString
```

5.4 Exercises

1. Write a Lean function that, given any element of *PropForm*, outputs a list of all the subformulas.

2. Write a Lean function that, given a list of propositional formulas and another propositional formula, determines whether the second is a logical consequence of the first.

3. Write a Lean function that, given a clause, tests whether any literal *Lit.pos p* appears together with its negation *Lit.neg p*. Write another Lean function that, given a formula in conjunctive normal form, deletes all these clauses.

DECISION PROCEDURES

We have seen that it is possible to determine whether or not a propositional formula is valid by writing out its entire truth table. This seems pretty inefficient; if a formula A has n variables, the truth table has 2^n lines, and hence checking it requires at least that many lines. It is still an open question, however, whether one can do substantially better. If $P \neq NP$, there is no polynomial algorithm to determine satisfiability (and hence validity.)

Nonetheless, there are procedures that seem to work better in practice. In fact, we can generally do *much* better in practice. In the next chapter, we will discuss *SAT solvers*, which are pieces of software that are are remarkably good at determining whether a propositional formula has a satisfying assignment.

Before 1990, most solvers allowed arbitrary propositional formulas as input. Most contemporary SAT solvers, however, are designed to determine the satisfiability of formulas in conjunctive normal form. Section 4.5 shows that, in principle, this does not sacrifice generality, because any propositional formula A can be transformed to an equivalent CNF formula B. The problem is that, in general, however, the smallest such B may be exponentially longer than A, which makes the transformation impractical. (See the exercises at the end of Chapter 4.) In the next section, we will show you an efficient method for associating a list of clauses to A with the property that A is satisfiable if and only if the list of clauses is. With this transformation, solvers can be used to test the satisfiability of any propositional formula.

It's easy to get confused. Remember that most formulas are neither valid nor unsatisfiable. In other words, most formulas are true for some assignments and false for others. So testing for validity and testing for satisfiability are two different things, and it is important to keep the distinction clear. But there is an important relationship between the two notions: a formula A is valid if and only if $\neg A$ is unsatisfiable. This provides a recipe for determining the validity of A, namely, use a SAT solver to determine whether $\neg A$ is satisfiable, and then change a "yes" answer to a "no" and vice-versa.

6.1 The Tseitin transformation

We have seen that if A is a propositional formula, the smallest CNF equivalent may be exponentially longer. The Tseitin transformation provides an elegant workaround: instead of looking for an *equivalent* formula B, we look for one that is *equisatisfiable*, which is to say, one that is satisfiable if and only if A is. For example, instead of distributing p across the conjunction in $p \vee (q \wedge r)$, we can introduce a new definition d for $q \wedge r$. We can express the equivalence $d \leftrightarrow (q \wedge r)$ in conjunctive normal form as

$$(\neg d \vee q) \wedge (\neg d \vee r) \wedge (\neg q \vee \neg r \vee d).$$

Assuming that equivalence holds, the original formula $p \vee (q \wedge r)$ is equivalent to $p \vee d$, which we can add to the conjunction above, to yield a CNF formula.

The resulting CNF formula implies $p \vee (q \wedge r)$, but not the other way around: $p \vee (q \wedge r)$ does not imply $d \leftrightarrow (q \wedge r)$. But the resulting formulas is equisatisfiable with the original one: given any truth assignment to the original one, we can give d the truth value of $q \wedge r$, and, conversely, for any truth assignment satisfying the resulting CNF formula, d *has* to have that value. So determining whether or not the resulting CNF formula is satisfiable is tantamount to determining whether the original one is. This may seem to be a roundabout translation, but the point is that the number of definitions is bounded

by the length of the original formula and the size of the CNF representation of each definition is bounded by a constant. So the length of the resulting formula is linear in the length of the original one.

The following code, found in the *LAMR* library in the *NnfForm* namespace, runs through a formula in negation normal form and produces a list of definitions *def_0*, *def_1*, *def_2*, and so on, each of which represents a conjunction or disjunction of two variables. We assume that none these *def* variables are found in the original formula. (A more sophisticated implementation would check the original formula and start the numbering high enough to avoid a clash.)

```
def defLit (n : Nat) := Lit.pos s!"def_{n}"

def mkDefs : NnfForm → Array NnfForm → Lit × Array NnfForm
  | lit l, defs    => (l, defs)
  | conj A B, defs =>
      let ⟨fA, defs1⟩ := mkDefs A defs
      let ⟨fB, defs2⟩ := mkDefs B defs1
      add_def conj (lit fA) (lit fB) defs2
  | disj A B, defs =>
      let ⟨fA, defs1⟩ := mkDefs A defs
      let ⟨fB, defs2⟩ := mkDefs B defs1
      add_def disj (lit fA) (lit fB) defs2
where
  add_def (op : NnfForm → NnfForm → NnfForm) (fA fB : NnfForm) (defs : Array␣
→NnfForm) :=
    match defs.findIdx? ((. == op fA fB)) with
    | some n => (defLit n, defs)
    | none   => let newdefs := defs.push (op fA fB)
                (defLit (newdefs.size - 1), newdefs)
```

The keyword *where* is used to define an auxiliary function that is not meant to be used anywhere else. The function *mkDefs* takes an NNF formula and a list of definitions, and it returns an augmented list of definitions and a literal representing the original formula. More precisely, the list *defs* is an array of disjunctions and conjunctions of variables, where the first one corresponds to *def_0*, the next one corresponds to *def_1*, and so on. In the cases where the original formula is a conjunction or a disjunction, the function first recursively creates definitions for the component formulas and then adds a new definition for the original formula. The auxiliary function *add_def* first checks to see whether the formula to be added is already found in the array. For example, when passed a conjunction $p \wedge def\_0$, if that formula is already in the list of definitions in position 7, *add_def* returns *def_7* as the definition of the formula and leaves the array unchanged.

To illustrate, we start by putting the formula

$$\neg(p \wedge q \leftrightarrow r) \wedge (s \rightarrow p \wedge t)$$

in negation normal form.

```
def ex1 := prop!{¬ (p ∧ q ↔ r) ∧ (s → p ∧ t)}.toNnfForm

#eval toString ex1
```

Removing extraneous parentheses, we get

$$((p \wedge q \wedge \neg r) \vee (r \wedge (\neg p \vee \neg q)) \wedge (\neg s \vee (p \wedge t)).$$

In the following, we compute the list of definitions corresponding to *ex1*, and then we use a little program to print them out in a more pleasant form.

```
#eval  ex1.mkDefs #[]

def printDefs (A : NnfForm) : IO Unit := do
```

```
  let ⟨fm, defs⟩ := A.mkDefs #[]
  IO.println s!"{fm}, where"
  for i in [:defs.size] do
    IO.println s!"def_{i} := {defs[i]}"

#eval printDefs ex1

/-
output:

def_7, where
def_0 := (p ∧ q)
def_1 := (def_0 ∧ (¬ r))
def_2 := ((¬ p) ∨ (¬ q))
def_3 := (r ∧ def_2)
def_4 := (def_1 ∨ def_3)
def_5 := (p ∧ t)
def_6 := ((¬ s) ∨ def_5)
def_7 := (def_4 ∧ def_6)
-/
```

We can obtain an equisatisfiable version of the formula by putting all the definitions into conjunctive normal form, collecting them all together, and adding one more conjunct with the variable d corresponding to the definition of the top level formula. There is a constant bound on the size of each CNF definition, which corresponds to a single conjunction or disjunction. (This would still be true if we added other binary connectives, like a bi-implication.) Since the number of definitions is linear in the size of the original formula, the size of the equisatisfiable CNF formula is linear in the original one.

There is an important optimization of the Tseitin transformation due to Plaisted and Greenbaum, who observed that for equisatisfiability, only one direction of the implication is needed for each subformula. If we start with a formula in negation normal form, each subformula occurs *positively*, which means that switching its truth value from negative to positive can only change the truth value of the entire formula in the same direction. In that case, only the forward direction of the implication is needed. For example, the formula $(p \wedge q) \vee r$ is equisatisfiable with $(d \vee r) \wedge (d \rightarrow p \wedge q)$, which can be expressed in CNF as $(d \vee r) \wedge (\neg d \vee p) \wedge (\neg d \vee q)$. To see that they are equisatisfiable, notice that any satisfying assignment to $(p \wedge q) \vee r$ can be extended to a satisfying assignment of $(d \vee r) \wedge (d \rightarrow p \wedge q)$ by giving d the same truth assignment as $p \wedge q$, and, conversely, $(d \vee r) \wedge (d \rightarrow p \wedge q)$ entails $(p \wedge q) \vee r$.

It isn't hard to put a implication of the form $(A \rightarrow B)$ into conjunctive normal form. The LAMR library defines a function to do that, and another one that turns the entire list of definitions returned by *mkDefs* into a single CNF formula.

```
def implToCnf (A B : NnfForm) : CnfForm :=
  (disj A.neg B).toCnfForm

def defsImplToCnf (defs : Array NnfForm) : CnfForm := aux defs.toList 0
  where aux : List NnfForm → Nat → CnfForm
  | [],         n => []
  | nnf :: nnfs, n => implToCnf (lit (defLit n)) nnf ++ aux nnfs (n + 1)
```

If we take the resulting formula and add a conjunct for the variable representing the top level, we have an equisatisfiable CNF formula, as desired.

A moment's reflection shows that we can do better. For example, if the formula A is already in CNF, we don't have to introduce any definitions at all. The following functions from the library do their best to interpret an NNF formula as a CNF formula, introducing definitions only when necessary. The first function, *NnfForm.orToCnf*, interprets a formula as a clause. For example, given the formula $p \vee (q \wedge \neg r) \vee \neg s$ and a list of definitions, it adds a definition d for $q \wedge \neg r$ and returns the clause $p \vee d \vee \neg s$. The function *NnfForm.andToCnf* does the analogous thing for conjunctions, and the function *NnfForm.toCnf* puts it all together.

```
def orToCnf : NnfForm → Clause → Array NnfForm → Clause × Array NnfForm
  | lit Lit.tr,  cls, defs  => ([Lit.tr], defs)
  | lit Lit.fls, cls, defs  => (cls, defs)
  | lit l, cls, defs        => (l :: cls, defs)
  | disj A B, cls, defs =>
      let ⟨cls1, defs1⟩ := orToCnf A cls defs
      let ⟨cls2, defs2⟩ := orToCnf B cls1 defs1
      (cls1.union cls2, defs2)
  | A, cls, defs =>
      let ⟨l, defs1⟩ := A.mkDefs defs
      (l::cls, defs1)

def andToCnf : NnfForm → Array NnfForm → CnfForm × Array NnfForm
  | conj A B, defs =>
    let ⟨fA, defs1⟩ := andToCnf A defs
    let ⟨fB, defs2⟩ := andToCnf B defs1
    (fA.union fB, defs2)
  | A, defs =>
    let ⟨cls, defs1⟩ := orToCnf A [] defs
    ([cls], defs1)

def toCnf (A : NnfForm) : CnfForm :=
  let ⟨cnf, defs⟩ := andToCnf A #[]
  cnf.union (defsImplToCnf defs)
```

The following example tests it out on *ex1.toCnf*. The comment afterward shows the resulting CNF formula and then reconstructs the definitions to show that the result is equivalent to the original formula.

```
#eval toString ex1.toCnf

/-
Here is ex1:

((p ∧ q ∧ ¬ r) ∨ (r ∧ (¬ p ∨ ¬ q)) ∧ (¬ s ∨ (p ∧ t))

Here is the CNF formula:

def_3 def_1,
def_4 -s,
-def_0 p,
-def_0 q,
-def_1 def_0,
-def_1 -r,
-def_2 -p -q,
-def_3 r,
-def_3 def_2,
-r -def_2 def_3,
-def_4 p,
-def_4 t,

Here we check to make sure it works:

def_0 := p ∧ q
def_1 := p ∧ q ∧ ¬ r
def_2 := ¬ p ∨ ¬ q
def_3 := r ∧ (¬ p ∨ ¬ q)
def_4 := p ∧ t
```

```
def_3 def_1 := (p ∧ q ∧ ¬ r) ∨ (p ∧ q ∧ ¬ r)
def_4 -s    := ¬ s ∨ (p ∧ t)

Each ':=' is really an implication.
-/
```

There is one additional optimization that we have not implemented in the library: we can be more efficient with iterated conjunctions and disjunctions. For example, you can check that $d \to \ell_1 \wedge \cdots \wedge \ell_n$ can be represented by the conjunction of the clauses $\neg d \vee \ell_1$ to $\neg d \vee \ell_n$, and $d \to \ell_1 \vee \cdots \vee \ell_n$ can be represented by the single clause $\neg d \vee \ell_1 \vee \cdots \vee \ell_n$.

6.2 Unit propagation and the pure literal rule

In earlier chapters, we considered only truth assignments that assign a truth value to all propositional variables. However, complete search methods for SAT like DPLL use *partial assignments*, which assign truth values to some subset of the variables. For a clause C and partial truth assignment τ, we denote by $[\![C]\!]_\tau$ the reduced clause constructed by removing falsified literals. If τ satisfies a literal in C, we interpret $[\![C]\!]_\tau$ as the singleton clause consisting of $\top$, while if τ falsifies all literals in C, then $[\![C]\!]_\tau$ is the empty clause, which we take to represent $\bot$. For a CNF formula Γ, we denote by $[\![\Gamma]\!]_\tau$ the conjunction of $[\![C]\!]_\tau$ with $C \in \Gamma$, where we throw away the clauses C such that $[\![C]\!]_\tau$ is $\top$. Remember that if $[\![\Gamma]\!]_\tau$ is an empty conjunction, we interpret it as being true. In other words, we identify the empty conjunction with $\top$. On the other hand, if $[\![\Gamma]\!]_\tau$ contains an empty clause, then $[\![\Gamma]\!]$ is equivalent to $\bot$.

Notice that we are reusing the notation $[\![A]\!]_\tau$ that we introduced in Section 4.2 to describe the evaluation of a formula A under a truth assignment τ. This is an abuse of notation, because in that section we were interpreting $[\![A]\!]_\tau$ as a *truth value*, whereas now we are taking $[\![C]\!]_\tau$ to be a *clause* and we are taking $[\![\Gamma]\!]_\tau$ to be *CNF formula*. But the abuse is harmless, and, indeed, quite helpful. Notice that if τ is a partial truth assignment that assigns values to all the variables in a clause C, then $[\![C]\!]_\tau$ evaluates to either the empty clause or $\top$, and if τ is a partial truth assignment that assigns values to all the variables in a CNF formula Γ, then $[\![\Gamma]\!]_\tau$ evaluates to either the empty CNF formula, $\top$, or the singleton conjunction of the empty clause, which is our CNF representation of $\bot$. In that sense, the notation we are using in this chapter is a generalization of the notation we used to describe the semantics of propositional logic.

Remember also that if τ is a truth assignment in the sense of Section 4.2, the semantic evaluation $[\![A]\!]_\tau$ only depends on the values of τ that are assigned to variables that occur in $[\![A]\!]$. In this chapter, when we say "truth assignment," we will generally mean "partial truth assignment," and the analogous fact holds: the values of $[\![C]\!]_\tau$ and $[\![\Gamma]\!]_\tau$ depend only on the variables found in C and Γ, respectively. In particular, if τ is a satisfying assignment for C, we can assume without loss of generality that τ assigns values only to the variables that occur in C, and similarly for Γ.

A key SAT-solving technique is *unit propagation*. Given a CNF formula Γ and a truth assignment τ, a clause $C \in \Gamma$ is *unit under* τ if τ falsifies all but one literal of C and the remaining literal is unassigned. In other words, C is unit under τ if $[\![C]\!]_\tau$ consists of a single literal. The only way to satisfy C is to assign that literal to true. Unit propagation iteratively extends τ by satisfying all unit clauses. This process continues until either no new unit clauses are generated by the extended τ or until the extended τ falsifies a clause in Γ.

For example, consider the partial truth assignment τ with $\tau(p_1) = \top$ and the following formula:

$$\Gamma_{\text{unit}} := (\neg p_1 \vee \neg p_3 \vee p_4) \wedge (\neg p_1 \vee \neg p_2 \vee p_3) \wedge (\neg p_1 \vee p_2) \wedge (p_1 \vee p_3 \vee p_6) \wedge$$
$$(\neg p_1 \vee p_4 \vee \neg p_5) \wedge (p_1 \neg p_6) \wedge (p_4 \vee p_5 \vee p_6) \wedge (p_5 \vee \neg p_6)$$

The clause $(\neg p_1 \vee p_2)$ is unit under τ because $[\![(\neg p_1 \vee p_2)]\!]_\tau = (p_2)$. Hence unit propagation will extend τ by assigning p_2 to $\top$. Under the extended τ, $(\neg p_1 \vee \neg p_2 \vee p_3)$ is unit, which will further extend τ by assigning p_3 to $\top$. Now the clause $(\neg p_1 \vee \neg p_3 \vee p_4)$ becomes unit and thus assigns p_4 to $\top$. Ultimately, no unit clauses remain and unit propagation terminates with $\tau(p_1) = \tau(p_2) = \tau(p_3) = \tau(p_4) = \top$.

Another important simplification technique is the *pure literal rule*. A literal l is called pure in a formula Γ if no clause in Γ contains the literal $\neg l$. The pure literal rule sets pure literals to $\top$.

In contrast to unit propagation, the pure literal rule can reduce the number of satisfying assignments. Consider for example for the formula $\Gamma = (p \vee q) \wedge (\neg q \vee r) \wedge (q \vee \neg r)$. The literal p is pure, so the pure literal rule will assign it to $\tau(p) = \top$. We leave it to you to check that Γ has three satisfying assignments, while $[\![\Gamma]\!]_\tau$ has only two satisfying assignments.

6.3 DPLL

One of the first and most well-known decision procedures for SAT problems is the Davis-Putnam-Logemann-Loveland (DPLL) algorithm, which was invented by Davis, Logemann, and Loveland in 1962, based on earlier work by Davis and Putnam in 1960. DPLL is a complete, backtracking-based search algorithm that builds a binary tree of truth assignments. For nearly four decades since its invention, almost all complete SAT solvers have been based on DPLL.

The DPLL procedure is a recursive algorithm that takes a CNF formula Γ and a truth assignment τ as input and returns a Boolean indicating whether Γ under the assignment τ can be satisfied. Each recursive call starts by extending τ using unit propagation and the pure literal rule. Afterwards, it checks whether the formula is satisfiable ($[\![\Gamma]\!]_\tau = \top$) or unsatisfiable ($[\![\Gamma]\!]_\tau = \bot$). In either of those cases it returns $\top$ or $\bot$, respectively. Otherwise, it selects an unassigned propositional variable p and recursively calls DPLL with one branch extending τ with $\tau(p) = \top$ and another branch extending τ with $\tau(p) = \bot$. If one of the calls returns $\top$ it returns $\top$, otherwise it returns $\bot$.

The effectiveness of DPLL depends strongly on the quality of the propositional variables that are selected for the recursive calls. The propositional variable selected for a recursive call is called a *decision variable*. A simple but effective heuristic is to pick the variable that occurs most frequently in the shortest clauses. This heuristics is called *MOMS* (Maximum Occurrence in clauses of Minimum Size). More expensive heuristics are based on *lookaheads*, that is, assigning a variable to a truth value and measuring how much the formula can be simplified as a result.

Consider the following CNF formula:

$$\Gamma_{\text{DPLL}} := (p_1 \vee p_2 \vee \neg p_3) \wedge (\neg p_1 \vee p_2 \vee p_3) \wedge (\neg p_1 \vee \neg p_2 \vee p_3) \wedge (p_1 \vee p_3) \wedge (\neg p_1 \vee \neg p_3)$$

Selecting p_3 as decision variable in the root node, results in a DPLL tree with two leaf nodes. However, when selecting p_2 as decision variable in the root node, the DPLL tree consists of four leaf nodes. The figure below shows the DPLL trees with green nodes denoting satisfying assignments, while black nodes denote falsifying assignments.

We have implemented a basic version of the DPLL procedure in Lean, which can be found in the *Examples* folder. We will start at the top level and work our way down. The function *dpllSat* accepts a CNF formula and returns either a satisfying assignment or *none* if there aren't any. It calls *propagateUnits* to perform unit propagation, and then calls *dpllSatAux* to iteratively pick an unassigned variable and split on that variable. We use the keyword *partial* to allow for an arbitrary recursive call.

```
partial def dpllSatAux (τ : PropAssignment) (φ : CnfForm) :
    Option (PropAssignment × CnfForm) :=
  if φ.hasEmpty then none
  else match pickSplit? φ with
  -- No variables left to split on, we found a solution.
  | none => some (τ, φ)
```

```
  -- Split on `x`.
  -- `<|>` is the "or else" operator, which tries one action and if that fails
  -- tries the other.
  | some x => goWithNew x τ φ <|> goWithNew (-x) τ φ

where
  /-- Assigns `x` to true and continues out DPLL. -/
  goWithNew (x : Lit) (τ : PropAssignment) (φ : CnfForm) :
      Option (PropAssignment × CnfForm) :=
    let (τ', φ') := propagateWithNew x τ φ
    dpllSatAux τ' φ'

/-- Solve `φ` using DPLL. Return a satisfying assignment if found, otherwise `none`.
↪-/
def dpllSat (φ : CnfForm) : Option PropAssignment :=
  let ⟨τ, φ⟩ := propagateUnits [] φ
  (dpllSatAux τ φ).map fun ⟨τ, _⟩ => τ
```

The function *pickSplit?* naively chooses the first variable it finds. It follows the Lean convention of using a question mark for a function that returns an element of an option type.

```
def pickSplit? : CnfForm → Option Lit
  | []      => none
  | c :: cs => match c with
    | x :: xs => x
    | _       => pickSplit? cs
```

The function *propagateUnits* performs unit propagation and returns the resulting formula and the augmented truth assignment. As long as there is a unit clause, it simplifies the formula and adds the new literal to the assignment.

```
partial def propagateUnits (τ : PropAssignment) (φ : CnfForm) : PropAssignment ×
↪CnfForm :=
  -- If `φ` is unsat, we're done.
  if φ.hasEmpty then ⟨τ, φ⟩
  else match φ.findUnit with
    -- If there are no unit clauses, we're done.
    | none   => ⟨τ, φ⟩
    | some x =>
      -- If there is a unit clause `x`, simplify the formula
      -- assuming `x` is true and continue propagating.
      let φ' := simplify x φ
      if τ.mem x.name
      then panic! s!"'{x}' has already been assigned and should not appear in the
↪formula."
      else propagateUnits (τ.withLit x) φ'
```

The function *propagateWithNew* is used to peform a split on the literal *x*. It assigns the value of *x* and then does unit propagation.

```
def propagateWithNew (x : Lit) (τ : PropAssignment) (φ : CnfForm) :
    PropAssignment × CnfForm :=
  propagateUnits (τ.withLit x) (simplify x φ)
```

Finally, the function *simplify* simplifies a CNF formula by assigning the literal *x* to true, assuming *x* is not one of the constants $\bot$ or ⊤.

```
def simplify (x : Lit) (φ : CnfForm) : CnfForm :=
  assert! x != lit!{⊥} && x != lit!{⊤}
  match φ with
  | []      => []
  | c :: cs =>
               let cs' := simplify x cs
               -- If clause became satisfied, erase it from the CNF
               if c.elem x then cs'
               -- Otherwise erase any falsified literals
               else c.eraseAll x.negate :: cs'

#eval toString <| simplify lit!{p} cnf!{p -q p -p, p, q q, -q -p, -p}
```

6.4 Autarkies and 2-SAT

Unit propagation and the pure literal rule form the core of the DPLL algorithm. In this section, we present a generalization of the pure literal rule, and we illustrate its use by providing an efficient algorithm for a restricted version of SAT.

The notion of an *autarky* or *autarky assignment* is a generalization of the notion of a pure literal. An autarky for a set of clauses is an assignment that satisfies all the clauses that are touched by the assignment, in other words, all the clauses that have at least one literal assigned. A satisfying assignment of a set of clauses is one example of an autarky. The assignment given by the pure literal rule is another. There are often more interesting autarkies that are somewhere in between the two.

Notice that if τ is an autarky for Γ, then $[\![\Gamma]\!]_\tau \subseteq \Gamma$. To see this, suppose C is any clause in Γ. If C is touched by τ, then $[\![C]\!]_\tau = \top$, and so it is removed from $[\![\Gamma]\!]_\tau$. If C is not touched by τ, then $[\![C]\!]_\tau = C$. Since every element of $[\![\Gamma]\!]_\tau$ is of one of these two forms, we have that $[\![\Gamma]\!]_\tau \subseteq \Gamma$.

Theorem

Let τ be an autarky for formula Γ. Then Γ and $[\![\Gamma]\!]_\tau$ are satisfiability equivalent.

Proof

If Γ is satisfiable, then since $[\![\Gamma]\!]_\tau \subseteq \Gamma$, we know that $[\![\Gamma]\!]_\tau$ is satisfiable as well. Conversely, suppose $[\![\Gamma]\!]_\tau$ is satisfiable and let τ_1 be an assignment that satisfies $[\![\Gamma]\!]_\tau$. We can assume that τ_1 only assigns values to the variables of $[\![\Gamma]\!]_\tau$, which are distinct from the variables of τ. Then the assignment τ_2 which is the union of τ and τ_1 satisfies Γ.

We now turn to a restricted version of SAT. A CNF formula is called k-SAT if the length of all clauses is at most k. 2-SAT formulas can be solved in polynomial time, while k-SAT is NP-complete for $k \geq 3$. A simple, efficient decision procedure for 2-SAT uses only unit propagation and autarky reasoning. The decision procedure is based on the following observation. Given a 2-SAT formula Γ that contains a propositional variable p, unit propagation on Γ using $\tau(p) = \top$ has two possible outcomes: (1) it results in a conflict, meaning that all satisfying assignments of Γ have to assign p to $\bot$, or (2) unit propagation does not result in a conflict, in which case the extended assignment after unit propagation is an autarky. To understand the latter case, note that assigning a literal in any clause of length two either satisfies the clause (if the literal is true) or reduces it to a unit clause (if the literal is false). So if there isn't a conflict, then it is impossible that unit propagation will produce a clause that is touched but not satisfied.

The decision procedure works as follows. Given a 2-SAT formula Γ, we pick a propositional variable p occurring in Γ and compute the result of unit propagation on Γ using $\tau(p) = \top$. If unit propagation does not result in a conflict, let τ' be the extended assignment and we continue with $[\![\Gamma]\!]_{\tau'}$. Otherwise let $\tau''(p) = \bot$ and we continue with $[\![\Gamma]\!]_{\tau''}$. This process is

repeated until the formula is empty, which indicates that the original formula is satisfiable, or contains the empty clause, which indicates that the original clause is unsatisfiable.

6.5 CDCL

The Conflict-Driven Clause-Learning (CDCL) decision procedure is the most successful SAT solving paradigm. Although it evolved from DPLL, modern implementations of CDCL have hardly any resemblance with the classic decision procedure. The algorithms differ in core algorithmic choices. When designing an algorithm, one typically needs to choose between making the algorithm smart or making it fast. Having it both ways is generally not an option due to conflicting requirements for the datastructures. State-of-the-art DPLL algorithms are designed to be smart: spend serious computational effort to pick the best variable in each recursive call to make the binary tree relatively small. As we will describe below, CDCL focuses on being fast. Unit propagation is by far the most computationally expensive step and everything in a CDCL solver is designed by make that as fast as possible. That prevents more complicated heuristics. Another important design decision for SAT solving algorithms is determining whether you plan to find a satisfying assignment or prove that none exists. DPLL focuses on finding a satisfying assignment by picking the most satisfiable branch first, while CDCL, as the name suggests, prefers conflicts and wants to find a short refutation. A satisfying assignment is a counterexample to no refutation exists.

The early development of CDCL is four decades later than DPLL. While DPLL simply backtracks when a conflict (leaf node) is reached, CDCL turns this conflict into a conflict clause that is added to the formula. The first and arguably most important step towards CDCL is the invention which clause to add: the first unique implication point. The first unique implication point is an invention by Marques-Silva and Sakallah (1996) and is still used in all modern day top-tier solvers.

The best decision heuristic for CDCL selects the unassigned variable that occurs most frequently in recently learned conflict clauses. This heuristic, called Variable-State Independent Decaying Sum (VSIDS in short), was introduced in the zChaff SAT solver by Moskewicz and co-authors (2001). The heuristic was been improved over the years, in particular by E'en and S"orensson (2003) in their solver MiniSAT. In recent years some other advances have been made as well. However, selecting variables in recent conflict clauses is still the foundation.

Adding clauses to the input formula could significantly increase the cost of unit propagation. Therefore clause learning was not immediately successful. The contribution that showed the effectiveness of clause learning in practice is the 2-literal watchpointers datastructure. In short, the solver does not maintain a full occurrence list of clauses, but only an occurrence list of the first two literals in a clause. These first two literals are required to have the following invariant: either both literals are unassigned or one of them is assigned to true. That invariant is enough to ensure that the clause is not unit. In case an assignment breaks the invariant, then the entire clause is explored to check whether it is possible to fix the invariant by swapping a satisfied or unassigned literal to the first two positions. If this is not possible, the clause is unit and the unassigned literal is assigned to true or the clause is falsified which triggers the learning procedure in the solver.

Another key difference between CDCL and DPLL is that the former restarts very frequently. Restarting sounds drastic, but simply means that all variables are unassigned. The remainder of the solver state, such as the heuristics and the learned clauses, are not altered. Modern SAT solver restart often: say roughly a thousand times per second (although this depends on the size of the formula). As a consequence, no binary search tree is constructed. Instead, CDCL might be best viewed a conflict-clause generator without any systematic search.

The are various other parts in CDCL solvers that are important for fast performance but that are beyond the scope of this course. For example, CDCL solvers do not only learn many clauses, but they also aggressively delete them to reduce the impact on unit propagation. Also, CDCL solvers use techniques to rewrite the formula, so called inprocessing techniques. In top-tier solvers, about half the runtime can be devoted to such techniques.

USING SAT SOLVERS

A satisfiability (SAT) solver determines whether a propositional formula has a satisfying assignment. The performance of SAT solvers has improved significantly in the last two decades. In the late 1990s, only formulas with thousands of variables and thousands of clauses could be solved. Today, many propositional formulas with millions of variables and millions of clauses can be solved. In this chapter, we will explain how to use SAT solvers and how to encode problems into propositional logic.

7.1 First examples

Remember that contemporary SAT solvers determine that satisfiability of propositional formulas in conjunctive normal form. Specifically, they use a format for specifying such formulas known as the *DIMACS* format. Our *LAMR* library proves a function that converts any CNF formula to that format, sends it to a SAT solver called *CaDiCaL*, and parses the answer. The following can be found in *Examples/using_sat_solver/examples.lean*:

```
def cadicalExample : IO Unit := do
  let (s, result) ← callCadical exCnf0
  IO.println "Output from CaDiCaL :\n"
  --IO.println s
  --IO.println "\n\n"
  IO.println (formatResult result)
  pure ()

#eval cadicalExample
```

It uses the same example CNF formulas defined in Section 5.3. You can change *exCnf0* to *exCnf1* or *exCnf2*, or use any other CNF formula you want. If you uncomment the two lines that begin *IO.println*, the Lean output will show you the raw output from CaDiCaL.

7.2 Encoding problems

All NP-complete problems can be transformed in polynomial time into a SAT problem (i.e., into a propositional formula). For many problems, such a transformation is quite natural. For some other problems, the transformation can be complicated. The transformation is not unique. Frequently there exist many way to encode a problem as a propositional formula. The encoding can have a big impact on the runtime of SAT solvers. Generally, the smallest encoding for a problem (in terms of the number of variables and the number of clauses) results in relatively strong performance. In this section we will describe a few encodings.

One way to encode a problem into propositional logic is to describe it first using some high-level constraints. Let's consider a couple of high-level constraints: Constrain a sequence of literals such that at least one of them is true (*atLeastOne*), or

that at most one of the is true (*atMostOne*), or that an odd number of them is true (*XOR*). Afterwards these constraints are encoded into propositional logical to obtain, so a SAT solver can be used to solve the resulting formula.

How to encode *atLeastOne*, *atMostOne*, and *XOR* as a set of clauses? The constraint *atLeastOne* is easy: simply use the disjunction of all the literals in the sequence. The second constraint is requires multiple clauses. The naive way generates a quadratic number of clauses: for each pair of literals (l_i, l_j) in the sequence, include the clause $\neg l_i \vee \neg l_j$. The naive way of encoding the *XOR* constraint results in an exponential number of clauses: all possible clauses over the literals such that an odd number of them are negated. For example, the encoding of $XOR(l_1, l_2, l_3)$ produces the following clauses: $l_1 \vee l_2 \vee \neg l_3, l_1 \vee \neg l_2 \vee l_3, \neg l_1 \vee l_2 \vee l_3, \neg l_1 \vee \neg l_2 \vee \neg l_3$

Although a quadratic number of clauses produced by can be acceptable *atMostOne* for a reasonable small sequence of literals, the exponential number of clauses produced by *XOR* would result in formulas that are hard to solve solely due to the size of the formula. Fortunately, one can encode both *atMostOne* and *XOR* using a linear number of clauses using the following trick: In case the sequence consists of more than four literals, split the constraint into two such that the first uses the first three literals of the sequence appended by a new literal y, while the second uses the remainder of the sequence appended by the literal $\neg y$. For example, *atMostOne*$(l_1, \ldots, l_n)$ is split into *atMostOne*(l_1, l_2, l_3, y) and *atMostOne*$(l_4, \ldots, l_n, \neg y)$. The splitting is repeated until none of the constraints has a sequence longer than four.

Another approach to encode a problem into propositional logic is to express it first as another NP-complete problem and afterwards transform the result into propositional logic. Let's demonstrate this approach for graph coloring. The graph coloring problem asks whether a given graph can be colored with a given number of colors such that adjacent vertices have different colors. Graph coloring problems can be easily encoded into a propositional formula, and SAT solvers can frequently solve these formulas efficiently.

Given a graph $G = (V, E)$ and k colors, the encoding uses $k|V|$ Boolean variables $x_{i,j}$ with $i \in \{1, \ldots, |V|\}$ and $j \in \{1, \ldots, k\}$ and $|V| + k|E|$ clauses. If $x_{i,j}$ is assigned to true it means that vertex i is assigned color j. The clauses encode two constraints: each vertex has a color and adjacent vertices have a different color. The first constraint can be encoded using a single clause per vertex. For example, for vertex i, we have the following clause: $x_{i,1} \vee \cdots \vee x_{i,k}$. The second constraint requires k binary clauses. For example, for an edge between vertices h and i, we have the following binary clauses: $(\neg x_{h,1} \vee \neg x_{i,1}) \wedge \cdots \wedge (\neg x_{h,k} \vee \neg x_{i,k})$.

The CNF formulas for a triangle (a fully connected graph with three vertices) for two colors and three colors is shown below. The first one is unsatisfiable, while the second one is satisfiable.

```
def triangleCnf2 := !cnf{
  x11 x12,
  x21 x22,
  x31 x32,
  -x11 -x21, -x12 -x22,
  -x11 -x31, -x12 -x32,
  -x21 -x31, -x22 -x32
}

def triangleCnf3 := !cnf{
  x11 x12 x13,
  x21 x22 x23,
  x31 x32 x33,
  -x11 -x21, -x12 -x22, -x13 -x23,
  -x11 -x31, -x12 -x32, -x13 -x33,
  -x21 -x31, -x22 -x32, -x23 -x23
}
```

Many problems, such as scheduling and planning, can naturally be expressed as a graph coloring problem. We can then transform the graph coloring problem into a SAT problem using the encoding described above.

An example of a problem that can be expressed as a graph coloring problem is the popular puzzle Sudoku: Place number is a grid consisting of nine squares subdivided into a further nine smaller squares in such a way that every number appears once in each horizontal line, vertical line, and square. This puzzle can be seen as a graph coloring where each small square

is a vertex, and vertices are connected if and only if the corresponding small squares occur the the same horizontal line, vertical line, or square. Below is one of the hardest sudoko puzzles with only 17 given numbers.

	4		3					
						7	9	
			6					
			1	4		5		
9							1	
2								6
			7	2				
	5					8		
			9					

If you take a look at the file *sudoku.lean* in the *Examples* folder, you will see that it is easily solved by a SAT solver.

7.3 Exercise: grid coloring

Ramsey Theory deals with patterns that cannot be avoided indefinitely. In this exercise we focus on a pattern of coloring a $n \times m$ grid with k colors: Consider all possible rectangles within the grid whose length and width are at least 2. Try to color the grid using k colors so that no such rectangle has the same color for its four corners. When this is possible, we say that the $n \times m$ grid is k-colorable while avoiding monochromatic rectangles. When using k colors, it is relatively easy to construct a valid $k^2 \times k^2$ grid. However, only few valid grids that are larger than $k^2 \times k^2$ are known. An example of a valid 3-coloring of the 9×9 grid is shown below.

```
0 0 1 1 2 2 0 1 2
2 0 0 1 1 2 2 0 1
1 2 0 0 1 1 2 2 0
0 1 2 0 0 1 1 2 2
2 0 1 2 0 0 1 1 2
2 2 0 1 2 0 0 1 1
1 2 2 0 1 2 0 0 1
1 1 2 2 0 1 2 0 0
0 1 1 2 2 0 1 2 0
```

Step 1. Encode whether there exists a coloring of the grid using three colors so that no such rectangle has the same color for its four corners. The encoding requires two types of constraints. First, each square needs to have at least one color. Second, if four squares form the corners of a rectangle, then they cannot have the same color.

Step 2. Solve the encoding for a 10×10 grid using a SAT solver and decode the solution into a valid coloring. Show the output of the SAT solver and a valid 3-coloring similar to the one above of the 9×9 grid.

Note that any valid coloring can be turned into another valid coloring by permuting the rows, columns, or colors. However, such valid colorings are isomorphic.

Step 3. Use the tool Shatter to (partially) break the symmetries of the encoding in Step 1 and count the number of solutions of the resulting formula.

7.4 Exercise: NumberMind

The game Number Mind is a variant of the well known game Master Mind.

Instead of colored pegs, you have to guess a secret sequence of digits. After each guess you're only told in how many places you've guessed the correct digit. So, if the sequence was 1234 and you guessed 2036, you'd be told that you have one correct digit; however, you would NOT be told that you also have another digit in the wrong place.

For instance, given the following guesses for a 5-digit secret sequence,

```
90342 ;2 correct
70794 ;0 correct
39458 ;2 correct
34109 ;1 correct
51545 ;2 correct
12531 ;1 correct
```

The correct sequence 39542 is unique.

Based on the following guesses,

```
5616185650518293 ;2 correct
3847439647293047 ;1 correct
5855462940810587 ;3 correct
9742855507068353 ;3 correct
4296849643607543 ;3 correct
3174248439465858 ;1 correct
4513559094146117 ;2 correct
7890971548908067 ;3 correct
8157356344118483 ;1 correct
2615250744386899 ;2 correct
8690095851526254 ;3 correct
6375711915077050 ;1 correct
6913859173121360 ;1 correct
6442889055042768 ;2 correct
2321386104303845 ;0 correct
2326509471271448 ;2 correct
5251583379644322 ;2 correct
1748270476758276 ;3 correct
4895722652190306 ;1 correct
3041631117224635 ;3 correct
1841236454324589 ;3 correct
2659862637316867 ;2 correct
```

Find the unique 16-digit secret sequence.

Step 1. Encode finding the correct sequence as a SAT problem. Use Boolean variables $x_{i,j}$ which are true if and only if at position i there is the digit j. The encoding consists of two parts: i) at each position there is exactly one digit; and ii) the correct number of digits from each line is matched. The encoding should *only* use these $x_{i,j}$ variables.

Step 2. Show that the correct sequence is unique. Which clause do you need to add to the encoding of Step 1?

Step 3. Reduce the size of the encoding by replacing the cardinality constraints from ii) in Step 1 by using the Sinz encoding.

DEDUCTION FOR PROPOSITIONAL LOGIC

In the study of computational complexity, a *language* is a set of strings over some alphabet. For example, we can consider the language *PROP* consisting of all propositional formulas, the language *SAT* consisting of all satisfiable formulas, and the language *TAUT* consisting of all tautologies. We have seen that *SAT* and *TAUT* are both *decidable*, which is to say, there is are algorithms to decide membership in those sets. The P=NP question is precisely the question as to whether there a polynomial time algorithm for *SAT*, or, equivalently, for *TAUT*.

Still speaking in broad terms, a *proof system* for a language is a relation $P(d, x)$ between strings with the property for any x, x is in the language if and only if there is a d such that $P(d, x)$ holds. In this case, we say that d is a *proof of membership* for x. We typically require that checking a proof is easy, say, by requiring that $P(d, x)$ runs in polynomial time. This is often easy to do by putting a lot of information into d. NP is the class of languages that have a polynomial time proof system with the additional property that for every x in the language, there is a proof of membership d whose length is polynomially bounded in the length of x. The language *SAT* is in NP because there are short proofs of satisfiability, namely, the satisfying assignments.

When it comes to propositional logic, when we talk about proof systems, we generally mean a proof system for *TAUT*. In other words, a proof system for propositional logic is supposed to show that a formula is valid. It can therefore also be used to establish unsatisfiability, since a formula A is unsatisfiable if and only if $\neg A$ is valid.

Assuming we have a particular proof system in mind, we write $\vdash A$ to mean that there is a proof of A. Remember that we use $\models A$ to mean that A is valid. The property that $\vdash A$ implies $\models A$ is known as *soundness*, and the property that $\models A$ implies $\vdash A$ is known as *completeness*. We want a proof system for propositional logic to be sound and complete.

Given that the set of tautologies in propositional logic is decidable, why do we need a proof system? The complexity of the decision procedures provides one answer: as far as we know, deciding whether or not something is a tautology takes exponential time in the worst case. From a theoretical standpoint, it is not clear whether proof systems can do substantially better; the question as to whether there is a polynomial time polynomially-bounded proof system for propositional logic is equivalent to the question as to whether NP = coNP, which is an open question. But, in practice, checking a proof is usually much more efficient than determining that something is a tautology from scratch.

Another concern is reliability. Fast decision procedures for propositional logic are highly optimized and sometimes buggy. Most modern SAT solvers can be asked to output a proof to justify the claim that the input is unsatisfiable. Checking the output with an independent checker adds confidence that the claim is correct.

Yet another reason to be interested in formal notions of proof is that they provide more faithful models of informal proof, the process by which mathematicians establish that mathematical claims are true. Finally, when we turn to first-order logic, we will see that there is no decision procedure for validity. In fact, even fairly restricted versions of the question can be undecidable. In cases like that, the best we can do is search for proofs and counterexamples, with no guarantee that either will succeed in finite time. In other words, proof systems for first-order logic are essential.

The notation $\Gamma \vdash A$ is used to express that A is provable from a set of hypotheses Γ. The notation $\vdash A$ therefore abbreviates $\emptyset \vdash A$. In this more general setting, soundness says that if $\Gamma \vdash A$, then $\Gamma \models A$, and completeness says that if $\Gamma \models A$, then $\Gamma \vdash A$. If Γ is the finite set $\{B_1, \ldots, B_n\}$, then $\Gamma \models A$ is equivalent to $\models B_1 \wedge \cdots \wedge B_n \rightarrow A$. So, for many purposes, we can focus on provability and validity without hypotheses. If the set Γ is infinite, however, we cannot express $\Gamma \models A$ in those terms. For most of this chapter, we will not worry about infinite sets of hypotheses, since

mechanized reasoning generally has to work with finite representations. But we will discuss the case where Γ is infinite in Section 8.5.

When it comes to talking about formal proofs, the words *proof*, *deduction*, and *derivation* are often used interchangeably. The last two are sometimes useful to distinguish formal derivations from ordinary (informal) mathematical proofs.

8.1 Axiomatic systems

Historically, one way of describing a proof system for propositional logic is to give a list of axioms, like this one:

1. $A \to (B \to A)$

2. $(A \to (B \to C)) \to ((A \to B) \to (A \to C))$

3. $A \to (B \to A \wedge B)$

4. $A \wedge B \to A$

5. $A \wedge B \to B$

6. $A \to A \vee B$

7. $B \to A \vee B$

8. $(A \to C) \to ((B \to C) \to (A \vee B \to C))$

9. $\neg\neg A \to A.$

These are really axiom *schemas*, which is to say, we have an axiom for every choice of A, B, and C. The only rule of inference in the system is *modus ponens*, which is the rule "from A and $A \to B$ conclude B." A *proof* of a formula A from hypotheses Γ is a sequence of formula $C_1, \ldots, C_m$ such that every C_i is either:

- an axiom,

- a hypothesis, or

- consequence of two earlier formulas C_j and C_k using modus ponens.

This proof system is sound and complete. Proving soundness is straightforward: you only need to check that each axiom is valid and that modus ponens preserves truth. This enables us to show, by induction, that each line of a proof $C_1, \ldots, C_m$ of A from Γ is true under an assignment τ, assuming every formula in Γ is. Proving completeness requires more work. But axiomatic systems are no longer of much practical interest: they do not provide convenient means of modeling informal proofs, and they are not useful for automated reasoning or search. We will therefore set them aside and focus our attention on other types of proof systems.

8.2 A sequent calculus

Let Γ be a finite set of propositional formula in negation normal form. The next calculus we will consider is designed to prove that the disjunction of the formulas in Γ is valid, which is to say, for every truth assignment τ, at least one of the formula in Γ is true. If A is a formula, we write Γ, A instead of $\Gamma \cup \{A\}$. The rules are as follows:

$$\frac{}{\Gamma, P, \neg P} \qquad \frac{\Gamma, A \qquad \Gamma, B}{\Gamma, A \wedge B} \qquad \frac{\Gamma, A, B}{\Gamma, A \vee B}$$

The first rule says that either something in Γ is true, or P is true, or $\neg P$ is true. The second rule says that if either something in Γ is true or A is true, and either something in Γ is true or B is true, then either something in Γ is true or $A \wedge B$ is true. The third rule says that if either something in Γ is true or A is true or B is true, then either something in Γ is true or $A \vee B$ is true. If we take these to be statements about truth values relative to some truth assignment τ, these rules are clearly sound with respect to the semantics.

The set Γ is called a *sequent*. (More specifically, it is called as *one-sided* sequent; we'll see two-sides sequents below.) A system of rules like this is therefore called a sequent calculus. As with axiomatic systems, we can think of a proof as a sequence of lines, but it is also common to represent proofs diagrammatically, as trees whose nodes are labeled by sequents. The following example provides a proof of the NNF equivalent of $A \wedge (B \vee C) \rightarrow (A \wedge B) \vee (A \wedge C)$.

$$
\cfrac{
  \cfrac{
    \neg A, \neg B \wedge \neg C, A \wedge B, A
    \qquad
    \cfrac{
      \neg A, \neg B \wedge \neg C, A, C
      \qquad
      \cfrac{
        \cfrac{
          \neg A, \neg B, B, C \qquad \neg A, \neg C, B, C
        }{
          \neg A, \neg B \wedge \neg C, B, C
        }
      }{
        \neg A, \neg B \wedge \neg C, A \wedge B, C
      }
    }{
      \neg A, \neg B \wedge \neg C, A \wedge B, A \wedge C
    }
  }{
    \cfrac{
      \cfrac{
        \neg A, \neg B \wedge \neg C, (A \wedge B) \vee (A \wedge C)
      }{
        \neg A \vee (\neg B \wedge \neg C), (A \wedge B) \vee (A \wedge C)
      }
    }{
      (\neg A \vee (\neg B \wedge \neg C)) \vee ((A \wedge B) \vee (A \wedge C))
    }
  }
}{}
$$

It is probably easiest to read this from the bottom up.

Remember that saying that a sequent Γ is *valid* in our semantics means that for every truth assignment τ, we have $[\![A]\!]_\tau = \top$ for some A in Γ.

Theorem

The sequent calculus presented above is sound and complete. In other words, a sequent Γ is provable if and only if it is valid.

Proof

The soundness direction is easy. Suppose there is a proof of Γ. Let τ be any truth assignment. We have already noted that each rule is sound, which is to say, if the premise or premises are true under an assignment, then so is the conclusion. By induction, we have that every sequent in the proof is true under assignment τ.

Proving completeness is usually trickier. In this case, we can use the fact that the rules of the calculus are bidirectional: any truth assignment that refutes the conclusion of a rule has to refute one of the premises. We can also use the fact that reading each rule backward decreases the number of binary connectives.

We prove the following by induction on the number of binary connectives: for every sequent Γ, either Γ is provable, or there is a truth assignment that makes every formula in Γ false. If Γ has no binary connectives, then it is a set of literals. If Γ contains a complementary pair of literals P and $\neg P$, it is an axiom, and otherwise there is an assignment that makes it false. For the inductive step, Γ must have a $\wedge$ or a $\vee$. Applying the corresponding rule, we have either the premises are valid or there is a counterexample to one of them. In the first case, Γ is valid, and in the second case, there is a counterexample to the conclusion.

Notice that this gives us yet another decision procedure for propositional logic: start with a sequent Γ (which can consist of a single formula, if you want), and apply the rules backward. If you reach an axiom on each leaf, you have a proof. If one branch fails to terminate with an axiom, reading off a counterexample to the leaf yields a counterexample to Γ. This is an important idea in automated reasoning, namely, it is desirable to search for a proof in such a way that failure implies the existence of a counterexample.

The sequent calculus is closely related to *tableau* proof systems that are also commonly used in automated reasoning. In the sequent proof above, there is a lot of needless repetition of formulas; tableau representations do a better job of recording only what changes as we go up the tree. Another difference is that tableau proof systems usually don't require

---

that the formulas are in negation normal form. Rather, the rules of a tableau system correspond to the sequent rules for the negation-normal-form equivalents. Since $A \rightarrow B$ is equivalent to $\neg A \lor B$, this requires changing A to $\neg A$ as we move up the tree. To avoid introducing new connectives, tableau systems often annotate formulas with emph{polarities}, so that A^+ represents A and A^- represents $\neg A$. The most jarring difference between sequent calculi and tableau systems is that the latter are often described in terms of search for a satisfying assignment rather than searching for a proof. For example, the rule in the sequent calculus that says

> to find a proof of $A \land B$, split and find a proof of A and a proof of B

becomes

> to find a satisfying assignment to $A \lor B$, split and find a satisfying assignment to A or a satisfying assignment to B.

The automated reasoning community is split between people who like to think in terms of searching for proofs and people who like to think in terms of searching for models. It is therefore important to learn how to speak both languages, and to be able to translate between them on the fly.

8.3 Resolution

Given that the sequent calculus implicitly corresponds to a decision procedure for propositional logic, it is natural to ask whether there is a proof system that corresponds more closely to DPLL, the decision procedure that was the focus of Chapter 6. We now describe such a system.

A *resolution* proof is designed to *refute* a CNF formula, that is, to prove that it is unsatisfiable. Let Γ be a CNF formula, represented as a set of clauses. As in Section 6.3, we can assume that none of the clauses contain repeated literals or a complementary pair, and we can think of each clause as a set of literals. If C is a clause and ℓ is a literal, we write C, ℓ for $C \lor \ell$. The *resolution rule* derives a new clause from an old one:

$$\frac{C, P \qquad C, \neg P}{C}$$

The rule says that if either C or P is true, and either C or $\neg P$ is true, then C has to be true. A *resolution proof* of a clause C from a set of clauses Γ is a sequence of steps (or a labelled tree) that obtains C from Γ using instances of the resolution rule. A *resolution refutation* of Γ is a resolution proof of the empty clause from Γ.

Theorem

A CNF formula Γ has a resolution refutation if and only if it is unsatisfiable.

Resolution can therefore be understood as a proof system for propositional logic in the following way: given any formula A, put $\neg A$ in CNF, and look for a resolution refutation. Such a refutation is a proof of A. The theorem above says that this system is sound and complete: A is valid if and only if $\neg A$ is unsatisfiable, which happens if and only if there is a refutation of $\neg A$.

Proof

Soundness follows straightforwardly from the fact that the resolution rule preserves truth under any truth assignment, while the empty clause is unsatisfiable.

To prove completeness, we use induction on the number of propositional variables to show that if Γ is unsatisfiable, there is a resolution refutation of Γ. If there are no variables, the fact that Γ is unsatisfiable means that it must be the set consisting of the empty clause, and we are done.

In the induction step, let P be any variable that occurs in Γ. If Γ is unsatisfiable, then so are $[\![\Gamma]\!]_{[P \mapsto \top]}$ and $[\![\Gamma]\!]_{[P \mapsto \bot]}$. By the inductive hypothesis, both of these are refutable.

Remember the relationship between $[\![\Gamma]\!]_{[P\mapsto\top]}$ and Γ: in the former, we remove all the clauses that include P and delete $\neg P$ from the remaining clauses. So a resolution refutation of the empty clause from $[\![\Gamma]\!]_{[P\mapsto\top]}$ uses only the clauses of $[\![\Gamma]\!]$ that don't contain P, possibly with $\neg P$ removed. Restoring $\neg P$ to all the initial clauses yields either a proof of the empty clause or a proof of $\neg P$.

In the first case, we have a proof of the empty clause from Γ, and we are done. Otherwise, applying the inductive hypotheses to $[\![\Gamma]\!]_{[\neg P\mapsto\bot]}$ and repeating the previous argument, we obtain either a proof of the empty clause or a proof of P. Once again, in the first case, we are done. Otherwise, we apply the resolution rule to the proof of P and the proof of $\neg P$, and we have a proof of the empty clause.

We can once again view the completeness proof as a decision procedure in disguise. In fact, the strategy of picking a variable and trying to refute $[\![\Gamma]\!]_{[P\mapsto\top]}$ and $[\![\Gamma]\!]_{[P\mapsto\bot]}$ simultaneously is exactly the splitting rule of DPLL, formulated in terms of demonstrating unsatisfiability rather than searching for a satisfying assignment. We can formulate the theorem above in more constructive terms as follows:

Theorem

For any CNF formula Γ, either Γ is satisfiable or there is a resolution refutation.

Remember that at any point in the DPLL search, we have a partial assignment τ that we are trying to extend to a satisfying assignment to $[\![\Gamma]\!]_\tau$. If we fail, we then have to give up on τ, backtrack, and try another path. To extract either a satisfying assignment or a resolution proof from the result, it suffices to show the following constructively:

> For any partial truth assignment τ, either $[\![\Gamma]\!]_\tau$ is satisfiable or there is a resolution proof of a clause C from Γ such that $[\![C]\!]_\tau = \bot$.

This yields the desired conclusion when τ is the empty assignment, since the only clause that evaluates to $\bot$ under the empty assignment is the empty clause.

We will sketch an explanation of how to read off the information above from the DPLL search, and we will leave it as an exercise for you to fill in the details. Remember that there are three steps that are interleaved in DPLL:

1. splitting on a variable p
2. unit propagation
3. removing pure literals

Unit propagation can be viewed as a special case of the splitting rule: If a CNF formula contains a unit clause with literal ℓ, then splitting on ℓ fails immediately on one branch, and the other branch corresponds to the result of applying unit propagation. Reasoning about the pure literal rule is more subtle. Suppose ℓ is pure in $[\![\Gamma]\!]_\tau$, and consider the DPLL search starting from $[\![\Gamma]\!]_{\tau[\ell\mapsto\top]}$, which is a subset of $[\![\Gamma]\!]_\tau$. If DPLL finds a satisfying assignment to $[\![\Gamma]\!]_{\tau[\ell\mapsto\top]}$, it can be extended to a satisfying assignment to $[\![\Gamma]\!]_\tau$ by mapping ℓ to $\top$. On the other hand, for any clause C such that $[\![C]\!]_{\tau[\ell\mapsto\top]} = \bot$, we have $[\![C]\!]_\tau = \bot$, because ℓ is pure in all the clauses of $[\![\Gamma]\!]_\tau$.

So we only have to deal with the splitting rule. A satisfying assignment for either branch results in a satisfying assignment for Γ, we only only have to prove the following:

> If there are a resolution proof of a clause C from Γ such that $[\![C]\!]_{\tau[P\mapsto\top]} = \bot$ and a resolution proof of a clause D from Γ such that $[\![D]\!]_{\tau[P\mapsto\bot]} = \bot$, then there is a resolution proof of a clause E from Γ such that $[\![E]\!]_\tau = \bot$.

We leave the proof of this fact to you.

8.4 Natural deduction

8.5 Compactness

USING LEAN AS A PROOF ASSISTANT

In Chapter 3, we considered the use of Lean as a programming language, and in Chapter 5 we saw that you can use Lean to define data types for things like propositional formulas and truth assignments, and thereby implement algorithms that act on these objects.

In this chapter, we will show how it is possible to represent propositional formulas directly in Lean's underlying foundation. In this sense, we are using Lean's foundation as an *object language* rather than a *metalanguage* for logic. To clarify the distinction, imagine implementing one programming language, such as Lisp, in another programming language, like C++. In this scenario, we can characterize Lisp as being the object language, that is, the one that is being implemented, and C++ as the metalanguage, the one that is carrying out the implementation. Defining propositional logic in Lean is similar: we are using one logical system, Lean, to implement another one, propositional logic. Our goal in this chapter is to clarify the sense in which Lean itself is a logical system, which is to say, its language can be used to state mathematical theorems and prove them.

9.1 Propositional Logic in Lean

9.2 Equational Reasoning in Lean

9.3 Structural Induction in Lean

A feature of working with a system like Lean, which is based on a formal logical foundation, is that you can not only define data types and functions but also prove things about them. The goal of this section is to give you a flavor of using Lean as a proof assistant. It isn't easy: Lean syntax is finicky and its error messages are often inscrutable. In class, we'll try to give you some pointers as to how to interact with Lean to construct proofs. The examples in this section will serve as a basis for discussion.

Remember that Lean's core library defines the *List* data type and notation for it. In the example below, we import the library, open the namespace, declare some variables, and try out the notation.

```
import Init

open List

variable {α : Type}
variable (as bs cs : List α)
variable (a b c : α)

#check a :: as
#check as ++ bs
```

```
example : [] ++ as = as := nil_append as

example : (a :: as) ++ bs = a :: (as ++ bs) := cons_append a as bs
```

The *variable* command does not do anything substantive. It tells Lean that when the corresponding identifiers are used in definitions and theorems that follow, they should be interpreted as arguments to those theorems and proofs, with the indicated types. The curly brackets around the declaration $\alpha : Type$ indicate that that argument is meant to be *implicit*, which is to say, users do not have to write it explicitly. Rather, Lean is expected to infer it from the context.

The library proves the theorems *[] ++ as* and *(a :: as) ++ bs = a :: (as ++ bs)* under the names *nil_append* and *cons_append*, respectively. You can see them by writing *#check nil_append* and *#check cons_append*. Remember that we took these to be the defining equations for the *append* function in Section 2.3. Although Lean uses a different definition of the append function, for illustrative purposes we will treat them as the defining equations and base our subsequent proofs on that.

Lean's library also proves *as ++ []* under the name *append_nil*, but to illustrate how proofs like this go, we will prove it again under the name *append_nil'*.

```
theorem append_nil' : as ++ [] = as := by
  induction as with
  | nil => rw [nil_append]
  | cons a as ih => rw [cons_append, ih]
```

In class, we will help you make sense of this. The *by* command tell Lean that we are going to write a *tactic* proof. In other words, instead of writing the proof as an expression, we are going to give Lean a list of instructions that tell it how to prove the theorem. At the start of the tactic proof, the theorem in question is our *goal*. At each step, tactics act on one more more of the remaining goals; when no more goals remain, the theorem is proved.

In this case, there are only two tactics that are needed. The *induction* tactic, as the name suggests, sets up a proof by induction, and the *rw* tactic *rewrites* the goal using given equations. Moving the cursor around in the editor windows shows you the goals at the corresponding state of the proof.

```
theorem append_assoc' : as ++ bs ++ cs = as ++ (bs ++ cs) := by
  induction as with
  | nil => rw [nil_append, nil_append]
  | cons a as ih => rw [cons_append, cons_append, ih, ←cons_append]
```

Here is a similar proof of the associativity of the *append* function. Note that the left arrow in the expression ←*cons_append* tell Lean that we want to use the equation from right to left instead of from left to right.

Now let us consider Lean's definition of the *reverse* function:

```
theorem reverse_def : reverse as = reverseAux as [] := rfl

theorem reverseAux_nil : reverseAux [] as = as := rfl

theorem reverseAux_cons : reverseAux (a :: as) bs = reverseAux as (a :: bs) := rfl
```

We will use these identities in the proofs that follow. Let's think about what it would take to prove the identity *reverse (as ++ bs) = reverse bs ++ reverse as*. Since *reverse* is defined in terms of *reverseAux*, we should expect to have to prove something about *reverseAux*. And since the identity mentions the *append* function, it is natural to try to characterize the way that *reverseAux* interacts with *append*. These are the two identities we need:

```
theorem reverseAux_append : reverseAux (as ++ bs) cs = reverseAux bs (reverseAux as␣
→cs) := by
  induction as generalizing cs with
```

```
  | nil => rw [nil_append, reverseAux_nil]
  | cons a as ih => rw [cons_append, reverseAux_cons, reverseAux_cons, ih]

theorem reverseAux_append' : reverseAux as (bs ++ cs) = reverseAux as bs ++ cs := by
  induction as generalizing bs with
  | nil => rw [reverseAux_nil, reverseAux_nil]
  | cons a as ih => rw [reverseAux_cons, reverseAux_cons, ←cons_append, ih]
```

Note the *generalizing* clause in the induction. What it means is that what we are proving by induction on *as* is that the identity holds *for every choice of bs*. This means that, when we apply the inductive hypothesis, we can apply it to any choice of the parameter *bs*. You should try deleting the *generalizing* clause to see what goes wrong when we omit it.

With those facts in hand, we have the identity we are after:

```
theorem reverse_append : reverse (as ++ bs) = reverse bs ++ reverse as := by
  rw [reverse_def, reverseAux_append, reverse_def, ←reverseAux_append', nil_append,
      reverse_def]
```

TERMS AND FORMULAS

ELEVEN

IMPLEMENTING TERMS AND FORMULAS

USING SMT SOLVERS

FIRST-ORDER LOGIC

FOURTEEN

IMPLEMENTING FIRST-ORDER LOGIC

FIFTEEN

USING FIRST-ORDER THEOREM PROVERS

SIXTEEN

SIMPLE TYPE THEORY